

Support at 1 billion weekly users

Lots of organizations deal with scale. Fewer deal with scale and hypergrowth.

The exercise starts with 150,000 support issues each week. People handle 60,000 of them.

My three-year planning scenario reaches 500,000 weekly issues while AI cuts the share handled by people from 40% today to 10% in Year 3.

1 billion weekly users	150k weekly issues	60% automated resolution	60k human-assisted issues
--	--	--	---

THREE-YEAR STRETCH CASE

~\$118M

If automation stays at today's 60% rate while demand grows, vendor support costs rise. The high case saves about \$138M in future vendor costs, or \$118M after \$20M in assumed rollout spending. The medium case saves about \$42M. [See the three outcomes.](#)

“Support has never really been about replying to just tickets. It’s about whether people get what they need, whether it actually serves them well.”

Glen Worthington, Head of User Ops

The three-year vision

AI transforms Support from a queue that responds to issues into a connected system that anticipates needs, resolves approved problems inside the product, and learns from every exception.


My recommendation

Plan toward 90% automated resolution by Year 3 as a stretch target, which would reduce the human share of weekly issues from 40% to 10%. Move trained people into complex cases, AI review, and repeat-issue fixes first. Reduce planned vendor seats once we reach our agreed Service Level Agreement (SLA) target and service coverage stays strong.

The starting figures come from the exercise. The three-year path is my planning scenario, separate from any OpenAI target.

AI handles routine issues while people take complex cases.

The bars show where weekly issues go as demand grows. In this stretch scenario, weekly users double and weekly issues per 10,000 users rise from 1.5 to 2.5 as paid, business, and developer use expands. AI handles more approved routine issues, while specialists focus on cases that need judgment.



CASES REACHING PEOPLE

40% → 10%

The human share falls across the plan. Weekly human cases rise in Year 1, then fall to 50k by Year 3.

CAPABILITY NEEDED

Interoperability across teams

Sales and Product Engagement bring customer context. Support Delivery needs deeper billing, account security, policy, and AI-quality skills, with Operations, Product, and Trust & Safety linked through one issue history.

OpenAI's [published support model](#) informs this vision. The growth and AI targets are my exercise assumptions.

Change the product, people, and process together.

First remove repeat issues at the source, then expand approved AI actions. Keep specialist coverage and quality checks in place as the customer mix changes.

PRODUCT Resolve more inside the product - Repeat-issue fixes from Product, Engineering, and Product Engagement - Approved in-product help and actions from AI PEOPLE Give people the harder work - Partners handle defined queues and review AI exceptions - Internal specialists own complex cases and AI quality - Trust & Safety reviews sensitive decisions PROCESS Run one learning loop - Operations links issue history across AI, partners, and specialists - One scorecard guides coaching, vendor decisions, and Product fixes			
First 90 days - Map eligible issues and baseline quality - Pilot invoice lookup and product guidance - Confirm owners with Product and Trust & Safety	Year 1 - Expand pilots after quality review - Build specialist and policy escalation playbooks - Protect coverage as demand rises	Year 2 - Extend proven AI paths across products - Use repeat failures to prioritize Product fixes - Cross-train partners for exceptions	Year 3 - Use approved AI actions across core journeys - Focus people on complex cases and AI review - Adjust vendor plans after SLA and coverage checks
Check every quarter Quality (Durable resolution, repeat contact, severe AI errors, tailored guidance) Productivity (Automated resolution, approved AI tool use, time to resolution) Capacity (Coverage breadth and depth, backlog age, attrition, no-shows) Cost to Serve (Cost per durable resolution, vendor spend, AI run cost, rework) Use one issue ID to compare AI, partners, and internal specialists on similar cases. During pilots, keep a small eligible group on the current route to measure AI's contribution. Sample AI closures to catch customers who leave without an answer.			

A \$118M high case, with clear funding gates.

Every case uses one demand forecast, with 220k, 330k, and 500k weekly issues across Years 1-3. Each compares vendor costs with the cost of keeping automation at today's 60% rate as demand grows. The net figures subtract assumed AI, tooling, training, and quality-check spending. Revenue growth stays outside this estimate.

LOW · SLOWER AUTOMATION -\$13.1M - Year 3 automation resolves 82% of weekly issues without a person - Only 10% of potential vendor savings reaches the budget - Full \$20M rollout spending exceeds savings	MEDIUM · PARTIAL VENDOR SAVINGS +\$42.4M - Year 3 automation resolves 90% of weekly issues without a person - Vendors would have handled 70% of shifted cases - Contract changes deliver part of the savings	HIGH · MORE VENDOR SAVINGS +\$117.5M - Year 3 automation resolves 90% of weekly issues without a person - Vendors would have handled 90% of shifted cases - Most planned vendor seat costs can be removed
--	---	--

The \$20M rollout figure is a planning estimate for AI use, tools, training, and quality checks across three years. It allocates \$4M to Year 1, \$6M to Year 2, and \$10M to Year 3. Finance would set actual budgets after pricing the pilots.

The model already uses 30% onshore and 70% offshore vendor costs . That mix is included in every case. [See the inputs and public price references.](#)

Separate location option	Moving 100 suitable vendor seats from onshore to offshore could save about \$7.0M per year before transition costs. That is separate from AI savings and carries potential quality risk in language, specialist accuracy, and escalation.
-------------------------------------	---

What I need from leadership Issue and cost data. Confirm which cases AI can resolve and which vendor costs can change. Named partners. Product, Engineering, Trust & Safety, Support Delivery, Operations, and Product Engagement can clear pilot decisions quickly. A capped pilot budget. Price two pilot journeys with Finance, then release further funding after results support expansion.	How I would judge progress Day 90. Target 10% fewer eligible cases reaching people, a 5-point gain in durable resolution, and zero severe AI errors in the two pilots. Each year. Check automation share against 70%, 82%, and 90%, alongside human share of 30%, 18%, and 10%. Each quarter. Require repeat contact and backlog age at or below baseline, plus vendor savings confirmed by Finance.
REWARD If quality holds and contracts adjust, the high case saves about \$118M net across three years.	RISK AND RESPONSE Wrong AI answers or rigid contracts can erase savings. Stop the affected AI path after a severe error, keep human coverage if queues rise, and release funding in stages.

How the three-year path was built.

PLANNING CASE	TODAY	YEAR 1	YEAR 2	YEAR 3
Weekly issues	150k	220k	330k	500k
Weekly users	1.0B	1.25B	1.5B	2.0B
Issues per 10k weekly users	1.5	1.76	2.2	2.5
Verified AI resolution	60%	70%	82%	90%
Human case share	40%	30%	18%	10%

Starting point. The exercise supplies 1 billion weekly users, about 150k weekly issues, and a 60% automated resolution rate.

Human case mix. As AI resolves more common issues, the remaining cases may need more judgment, current guidance, and stronger escalation paths.


AI-ready issues. Before expanding AI, I would confirm the volume of issues it can handle safely with approved answers or actions, then check that customers stay helped.

Demand growth. This stretch case assumes weekly users double over three years. Paid, business, and developer use can also raise issue frequency from 1.5 to 2.5 per 10,000 weekly users. Together those assumptions produce 500k weekly issues in Year 3. The yearly steps are planning estimates, and I would replace them with observed segment demand.

Slower AI progress. At 82% AI resolution in Year 3, about 90k weekly cases still reach people. I would keep coverage until the issue mix and service levels support a change.

Staffing. This scenario sets case-share targets and leaves headcount open. Vendor plans follow observed demand, quality, coverage, and contract terms.

One issue history from start to finish.



Inputs behind the financial range.

INPUT	HIGH CASE	REASON
Vendor rep cost	\$93.6k onshore \$23.4k offshore	Public prices of \$45 hourly onshore and \$1,950 monthly offshore.
Vendor location mix	30% onshore 70% offshore	Already included in the \$44.46k yearly blended price in every case.
Cases vendors would handle	90%	AI takes common cases that partners may otherwise handle. Actual issue ownership must confirm this.
Cases per vendor rep	3,000 yearly	About 12 cases per scheduled day. Case logs must confirm handling time.
Savings that reach the budget	40%, 75%, 90%	Contract changes and fewer planned seats take effect over Years 1-3.
AI and rollout spending	\$4M, \$6M, \$10M	Illustrative Year 1-3 costs for AI use, tools, training, and quality checks. Finance would price the pilots first.

The high case saves \$137.5M in vendor costs against the flat-60% AI forecast and spends \$20M on AI and rollout, leaving \$117.5M in modeled savings. Today's vendor bill could still rise as total issue volume grows.

Two checks on the high case. If a vendor rep handles 6,000 cases yearly instead of 3,000, modeled savings fall to \$48.8M after rollout spending. If automation reaches 65%, 68%, and 70% in Years 1-3 without this plan, modeled savings fall to \$69.5M. New engineering salaries would need a separate budget estimate.

Public vendor price references [SupportOps Global](#), [ConnectPro](#), and [PITON-Global](#). OpenAI contract prices may differ from these planning references.