

Support at 1 billion weekly users

Lots of organizations deal with scale. Fewer deal with scale and hypergrowth.

The exercise starts with 150,000 support issues each week. People handle 60,000 of them.

My three-year planning scenario reaches 500,000 weekly issues while AI cuts the share handled by people from 40% today to 10% in Year 3.

Imagine a World where

2 billion weekly users	500k weekly issues	90% automated resolution	50k human-assisted issues
--	--	--	---

The three-year vision

AI transforms Support from a queue that just responds to issues, to a connected system that 1) anticipates needs, 2) resolves approved problems inside the product, and 3) learns from every exception.

My recommendation

1. Plan toward 90% automated resolution by Year 3 as a stretch target, reducing the human share of weekly issues from 40% to 10%.


2. Move trained people into complex cases, AI review, and repeat-issue fixes first.

3. Reduce planned vendor seats once we reach our agreed Service Level Agreement (SLA) target and service coverage stays strong.

The starting figures come from the exercise. The three-year path is my planning scenario, separate from any OpenAI target.

AI handles routine issues while people take complex cases.

The bars show where weekly issues go as demand grows. In this stretch scenario, weekly users double and weekly issues per 10,000 users rise from 1.5 to 2.5 as paid, business, and developer use expands. AI handles more approved routine issues, while specialists focus on cases that need judgment.



CASES REACHING PEOPLE

40% → 10%

- Human cases rise from 60k to 66k weekly in Year 1 because total demand grows faster than early automation gains.
- They stay near 66k in Year 2, then fall to 50k in Year 3 as AI handles more approved tasks and Product fixes repeat issues.

CAPABILITY NEEDED

Interoperability across teams

- Sales and Product Engagement bring customer context.
- Support Delivery builds deeper billing, account security, policy, and AI-quality skills while Operations, Product, and Trust & Safety share one issue history.

OpenAI's [published support model](#) informs this vision. The growth and AI targets are my exercise assumptions.

Change the product, people, and process together.

First remove repeat issues at the source, then expand approved AI actions. Keep specialist coverage and quality checks in place as the customer mix changes.

PRODUCT

Resolve more inside the product

- Repeat-issue fixes from Product, Engineering, and Product Engagement
- AI answers from approved guidance and completes bounded tasks, such as invoice lookup, with a human handoff when needed

PEOPLE

Give people the harder work

- Partners handle defined queues and review AI exceptions
- Specialists own complex cases and AI quality
- Trust & Safety reviews sensitive decisions

PROCESS

Run one learning loop

- Operations links issue history across AI, partners, and specialists
- One scorecard guides coaching, vendor decisions, and Product fixes

First 90 days

- Audit:** Baseline demand, friction, vendors, staffing, tools, and skills
- Align:** Owners, scorecards, and one urgent service gap
- Scale:** Pilot invoice lookup and product guidance, coach adoption, and set next-quarter priorities

Year 1

- Put approved AI help inside high-volume product journeys
- Link Support Delivery, Product, and Trust & Safety decisions
- Plan about 1,030 modeled vendor seats as demand grows

Year 2

- Expand approved AI actions across products and channels
- Route repeat failures into Product fixes and partner coaching
- Plan about 1,030 modeled vendor seats while AI keeps pace with growth

Year 3

- Run AI as the first route for eligible support needs
- Keep specialists on complex cases, policy, and AI quality
- Plan about 780 vendor seats after SLA and coverage checks

Why those pilots?

Invoice lookup and product guidance have clear answers and human escalation. The first audit would confirm demand and risk before launch.

How the seats were modeled.

Year 3 uses 500k weekly issues × 10% reaching people × 52 weeks × 90% vendor share ÷ 3,000 cases per rep yearly = 780 seats. This is a scenario estimate, not OpenAI staffing.

● Start

Expand approved AI paths when customer results hold.

● Stop

Pause an affected action after a severe error.

● Continue

Keep human coverage while case quality needs it.

Check every quarter

Quality

Durable resolution, repeat contact, severe AI errors, tailored guidance

In quarterly top-to-top reviews, OpenAI and each vendor's leadership would assess these four areas against shared targets. For any gap, I would agree on a vendor-specific plan with an owner, milestones, and a follow-up date.

Productivity

Automated resolution, approved AI tool use, time to resolution

Capacity

Coverage breadth and depth, backlog age, attrition, no-shows

Cost to Serve

Cost per durable resolution, vendor spend, AI run cost, rework

A \$114M high case, with clear conditions.

Every case uses one demand forecast, with 220k, 330k, and 500k weekly issues across Years 1-3. Each compares vendor costs with the cost of keeping automation at today's 60% rate as demand grows. The net figures subtract assumed AI, tooling, training, and quality-check spending. Revenue growth stays outside this estimate.

LOW · SLOWER AUTOMATION -\$13.5M - AI progress stalls at the Year 2 target of 80% in Year 3 - About \$6.5M in vendor costs saved \$20M in assumed added spending produces a net loss	MEDIUM · PARTIAL VENDOR SAVINGS +\$41.1M - AI resolves 90% of weekly issues by Year 3 - Vendor teams handle fewer cases as AI expands - About \$61.1M saved before added spending	HIGH · MORE VENDOR SAVINGS +\$114.1M - AI resolves 90% of weekly issues by Year 3 - Modeled vendor capacity falls to about 780 seats - About \$134.1M saved before added spending
---	--	--

I assumed \$20M in additional AI use, tooling, training, and quality checks over three years, with \$4M, \$6M, and \$10M in Years 1-3. This is a round planning input, not a vendor quote or an OpenAI budget. Finance would price the two 90-day pilots before setting a budget.

The model already uses 30% onshore and 70% offshore vendor costs . That mix is included in every case. [See the inputs and public price references.](#)

Separate location option	Moving 100 suitable vendor seats from onshore to offshore could save about \$7.0M yearly before transition costs. That is separate from AI savings. The high case starts with about 940 modeled vendor seats, and offshore shifts carry quality risk in language, specialist accuracy, and escalation.
-------------------------------------	--

What I need from leadership Two pilot journeys. Approve invoice lookup and product how-to guidance as the first low-risk AI tests, with human escalation available. Interoperability alignment. Assign Product, Engineering, Trust & Safety, Support Delivery, Operations, and Product Engagement to one issue history and fast decisions. A priced pilot budget. Ask Finance to cost the two pilots so AI usage and quality checks have a clear limit.	How I would judge progress Day 90. Target 10% fewer eligible cases reaching people, a 5-point gain in durable resolution, and zero severe AI errors in the two pilots. Each year. Check automation share against 70%, 80%, and 90%, alongside human share of 30%, 20%, and 10%. Each quarter. Require repeat contact and backlog age at or below baseline, plus vendor savings confirmed by Finance.
--	---

“Support has never really been about replying to just tickets. It’s about whether people get what they need, whether it actually serves them well.”

Glen Worthington, Head of User Ops

How the three-year path was built.

PLANNING CASE	TODAY	YEAR 1	YEAR 2	YEAR 3
Weekly issues	150k	220k	330k	500k
Weekly users	1.0B	1.25B	1.5B	2.0B
Issues per 10k weekly users	1.5	1.76	2.2	2.5
Verified AI resolution	60%	70%	80%	90%
Human case share	40%	30%	20%	10%

Starting point. The exercise supplies 1 billion weekly users, about 150k weekly issues, and a 60% automated resolution rate.

Human case mix. As AI resolves more common issues, the remaining cases may need more judgment, current guidance, and stronger escalation paths.


AI-ready issues. Before expanding AI, I would confirm the volume of issues it can handle safely with approved answers or actions, then check that customers stay helped.

Demand growth. This stretch case assumes weekly users double over three years. Paid, business, and developer use can also raise issue frequency from 1.5 to 2.5 per 10,000 weekly users. Together those assumptions produce 500k weekly issues in Year 3. The yearly steps are planning estimates, and I would replace them with observed segment demand.

Slower AI progress. The low case holds Year 3 AI resolution at the Year 2 stretch target of 80%, leaving about 100k weekly cases for people. This is an illustrative downside, not a forecast.

Staffing. Modeled vendor seats rise from about 940 today to 1,030 in Year 1, stay near 1,030 in Year 2, then fall to 780 in Year 3. These are workload estimates, not current OpenAI headcount. Actual staffing depends on case mix, SLA, quality, and existing vendor arrangements.

One issue history from start to finish.



Inputs behind the financial range.

INPUT	HIGH CASE	REASON
Vendor rep cost	\$93.6k onshore \$23.4k offshore	Public prices of \$45 hourly onshore and \$1,950 monthly offshore.
Vendor location mix	30% onshore 70% offshore	Already included in the \$44.46k yearly blended price in every case.
Human cases handled by vendors	90%	An illustrative staffing assumption. Medium and low cases use 70%. Actual case ownership needs review.
Cases per vendor rep	3,000 yearly	About 12 cases per scheduled day. Case logs must confirm handling time.
Potential savings in the budget	40%, 75%, 90%	The modeled share rises as future staffing plans change. Existing vendor terms must allow those reductions.
Additional AI spending	\$4M, \$6M, \$10M	A \$20M three-year planning input, not a quote. Finance would price the pilots first.

The high case saves \$134.1M in vendor costs against the flat-60% AI forecast and spends an assumed \$20M on AI and rollout, leaving \$114.1M in modeled savings. The \$20M is a round input, equivalent to about \$1.62 per additional AI-resolved issue across the three-year plan. It comes from no vendor quote. Today's vendor bill could still rise as total issue volume grows.

Two checks on the high case. If a vendor rep handles 6,000 cases yearly instead of 3,000, modeled savings fall to \$47.0M after added spending. If automation reaches 65%, 68%, and 70% in Years 1–3 without this plan, modeled savings fall to \$66.1M. New engineering salaries would need a separate budget estimate.

Public vendor price references [SupportOps Global](#), [ConnectPro](#), and [PITON-Global](#). OpenAI contract prices may differ from these planning references.