

Support at one billion weekly users

I would first verify how many customer issues stay resolved. Then I would plan the human capacity that remains and improve quality, productivity, capacity, and cost as four linked operating levers. Those are the lenses I used to manage vendor programs at Meta, adapted here for Support.

1B weekly users	150k weekly issues	60% automated resolution	60k human-assisted issues
---------------------------	------------------------------	------------------------------------	-------------------------------------

My first decision

I would ask for 90 days to audit the baseline and test two high-volume, low-risk journeys. I would keep trained partner and specialist capacity in place while those pilots produce evidence.

I managed five BPO partners across onshore and offshore Meta programs for four years. That was sales operations, so I would change the measures before applying the management approach to Support. At Rowland AI, I build agent workflows with evaluations and human review, which informs the automation pilots here.

My planning case reaches about 281k weekly issues.

The exercise gives us 150k issues each week and a 60% automated-resolution rate. I assumed weekly users grow 1.5× over three years, issues per user rise 1.25×, and the average human case takes 1.3× today's effort as products and customer needs expand.

PLANNING VIEW	WEEKLY ISSUES	HUMAN CASES	HUMAN EFFORT UNITS
Today, exercise baseline 60% automated	150k	60k	60k
Year 3, current rate holds 60% automated	281k	112.5k	146k
Year 3, tested improvement 75% verified automation	281k	70.3k	91k

One effort unit equals the average effort for one human-assisted issue today. I chose 75% as a planning case to test. The exercise gives no automation target.

Capacity decision

Even the improved case raises human effort about 52% above today's baseline. I would maintain a flexible partner reserve and use observed handle time, occupancy, shrinkage, service windows, and case mix before converting this model into staffing.

I would run one scorecard across four levers.

Each measure should answer a management question and prompt a decision. I would compare partners on similar issue and risk mixes. One linked issue ID would show the customer outcome alongside each automation, partner, and specialist step that contributed to it.

01	Quality	MEASURE Audited answer accuracy, seven-day repeat contact, and severe errors. See a fictional quality case ↗	DECISION Pause a route after a severe error. Calibrate scoring and fix the content or handoff before expansion.
02	Productivity	MEASURE Durable resolutions per paid hour, rework, and transfers by issue type. See a fictional coverage case ↗	DECISION Check the denominator, then improve tools, routing, or coaching where useful capacity falls.
03	Capacity	MEASURE Arrivals against qualified production hours, backlog age, attendance, and attrition. See a fictional staffing case ↗	DECISION Add trained flex or change schedules when the forecast and roster diverge.
04	Cost	MEASURE Total AI, partner, specialist, quality, and rework cost per durable resolution. See a fictional vendor case ↗	DECISION Compare matched cases before shifting volume toward a lower-price partner.

The linked demonstrations use fictional sales and vendor records. They show how I would diagnose a gap and assign a recovery plan. The support measures above replace the sales metrics in those examples.

I would keep policy and judgment inside OpenAI.

One issue record should follow the customer across self-service, automation, a partner, and an internal specialist. That record needs the issue type, prior attempts, risk level, handoffs, and final outcome.

Product help and AI Handle tested, low-risk requests with a clear route to a person when the answer fails or the issue carries risk.	Partner teams Handle repeatable cases in trained queues, document failures, and send patterns back to the knowledge and product owners.	Internal specialists Own sensitive cases, unclear policy, product defects, incidents, and the final decision on escalations.
--	---	--

OpenAI would own

Routing rules, knowledge, AI release gates, forecasts, quality calibration, partner contracts, and Product and Engineering priorities.

Partners would own

Staffing, coaching, and execution in approved queues, with audited outcomes and a tested backup route for peaks.

Expected effect and tradeoff

I expect better product help to reduce repeatable demand, partner coaching to improve durable resolution, and trained flex capacity to limit backlog swings. I would accept slower automation expansion and some reserve cost until the customer outcomes support more volume.

I would review service and staffing weekly, calibrate quality across sites monthly, and revisit contracts quarterly. I would use a base capacity commitment with a reward tied to audited quality and resolution. A cheaper unit price only helps when repeat contact and specialist rework remain healthy.

I would scale the journeys that pass the review.

First 90 days	Audit automated resolution and recontact. Join issue records across channels. Establish the four-part scorecard and a weekly demand forecast. Choose two low-risk journeys for controlled pilots.
Year 1	Test product help, agent assistance, and partner routing against comparison groups. Add specialist escalation and incident playbooks before increasing volume.
Years 2–3	Expand the journeys that improve durable resolution, train backup partner capacity, and shift recurring issue clusters into Product and Engineering fixes.

Pause A severe privacy, fraud, policy, or account-access error pauses the affected route and sends the issue to a human owner.	Investigate Two weekly cohorts where recontact exceeds the matched comparison or accuracy trails it stop expansion while the team reviews cases.	Expand Durable resolution improves while severe errors and backlog stay stable, so the team can add volume in stages.
--	--	---

My ask of leadership

Give Support Operations access to linked issue and outcome data, two pilot journeys, a flexible partner reserve, and a Product and Engineering review partner. I would return after 90 days with a measured staffing plan and a decision on each pilot.

These are the assumptions I would test first.

Planning assumptions

The exercise supplies 1B weekly users, 150k weekly issues, and 60% automated resolution. My year-three case assumes 1.5× user growth, 1.25× more issues per user, and 1.3× human effort per case. A 75% verified automated rate is a conditional planning case.

At 60% automation, $281,250 \times 40\% \times 1.3$ gives 146,250 effort units. At 75%, $281,250 \times 25\% \times 1.3$ gives 91,406 units. A stress case with 65% automation and 1.5× human effort reaches about 148k units.

Seven-day repeat contact is a starting proxy for durable resolution. Longer issue types need a longer review window. I would validate issue mix, handling time, and current partner capacity before setting a hiring number.

External evidence

- [OpenAI's published support model](#) connects customer surfaces, knowledge, evaluation, and frontline feedback.
- [COPC CX Standard](#) informs customer operations and partner management.
- [ISO 18295-2](#) covers client oversight of outsourced and internal contact centres.
- [Generative AI at Work](#) studied AI assistance for 5,179 support agents. I would measure local effects before using an efficiency gain in the forecast.
- [NIST AI Risk Management Framework](#) informs monitoring, human review, and recovery.

This is my independent recommendation for a hypothetical exercise. The scenarios use sample records and goals. I have no access to OpenAI internal support data.