

Support at 1 billion weekly users

Lots of organizations deal with scale. Fewer deal with scale and hypergrowth. The exercise already starts with 150,000 issues and 60,000 human-assisted cases each week. A small increase in issue frequency creates a large amount of additional demand. In my three-year planning case, issues approach 281,000 each week, while a harder human queue requires 52% more effort even with stronger automation. Without a change in the operating model, slower answers and repeat contacts become a bigger risk.

1 billion weekly users	150k weekly issues	60% automated resolution	60k human-assisted issues
--	--	--	---

“Support has never really been about replying to just tickets. It’s about whether people get what they need, whether it actually serves them well.”

Glen Worthington, Head of User Ops

The vision

What this looks like in 3 years

AI resolves approved routine issues inside the product and gives skilled people the history and guidance they need for harder cases. Customers can move between AI, partners, and internal specialists without starting over. Every interaction improves knowledge, evaluations, and product decisions, while people retain judgment on sensitive issues.

How I approached this assignment

- Diagnose demand and case complexity
- Design the operating model
- Deliver in measured stages

The starting figures come from the hypothetical take-home exercise. The three-year figures are my planning assumptions.

Weekly demand nearly doubles; human effort still rises.

I use a three-year planning case to show the capacity implications. By Year 3, weekly users reach 1.5× today's exercise baseline, issues per user reach 1.25×, and a human case takes 1.3× today's average effort. Verified automated resolution rises from 60% to 75% in stages.



PLANNING VIEW	WEEKLY ISSUES	VERIFIED AUTOMATION	HUMAN CASES	HUMAN EFFORT UNITS
Today 1.00× users · 1.00× issues/user · 1.00× effort	150k	60%	60k	60k
Year 1 1.15× users · 1.08× issues/user · 1.10× effort	186k	65%	65k	72k
Year 2 1.32× users · 1.16× issues/user · 1.20× effort	230k	70%	69k	83k
Year 3 1.50× users · 1.25× issues/user · 1.30× effort	281k	75%	70k	91k

Why these assumptions A larger user base raises demand. More products and business use cases can lift issue frequency. Automation can remove easier cases first, leaving a harder human queue. The multipliers illustrate those pressures; they are my judgment, separate from any OpenAI forecast.

One effort unit equals today's average human case. The figures are rounded estimates. If automation stays at 60% in Year 3, the model reaches about 113k human cases and 146k effort units.

What changes for capacity and capability

Case count grows slowly only if more issues stay resolved through automation. Human effort still rises about 52%, so I would add specialized skills, better tools, and trained flexible coverage. Staffing decisions would use handle time by case type, qualified hours, occupancy, shrinkage, and service windows.

AI-led Support that learns from every interaction.

[OpenAI's published support model](#) connects customer surfaces, living knowledge, evaluations, and frontline feedback. I would extend that approach across the next three years. AI should resolve approved routine issues, help people handle complex ones, and turn repeated failures into better guidance and product changes.

The strategy

Where to play and how to win

Where to play

Start with common, low-risk issues where AI can provide an answer or approved action and the outcome can be checked. Keep sensitive, unclear, and high-risk issues with trained people.

How to win

Connect every customer step to one issue history. Give AI and people current knowledge, evaluate answers against customer outcomes, and send repeat failures to Product and Engineering.

01 / CUSTOMER HELP

Resolve routine issues in product

Invest in guided help and approved AI actions with a clear path to a person.

Expected benefit

Faster answers and less repeatable demand.

Trade-off

Start with a narrow set of low-risk journeys.

02 / LEARNING LOOP

Improve each answer

Join issue history, living knowledge, evaluations, and frontline feedback.

Expected benefit

Better quality across channels and fewer recurring failures.

Trade-off

Data and quality setup come before broad rollout.

03 / HUMAN CAPACITY

Build for complex cases

Train partners for defined queues and internal specialists for judgment and product defects.

Expected benefit

Reliable coverage as the remaining case mix gets harder.

Trade-off

Training and backup coverage add cost early.

This order protects quality while demand grows. A higher automation rate helps only when customers stay resolved and severe errors remain controlled.

One customer issue, one history, clear owners.

The operating model begins with one issue record. It holds the customer's question, previous attempts, risk level, handoffs, and final outcome. Routing then sends the issue to the right capability instead of forcing every request through one queue.

01 Capture the issue Collect context and prior attempts once.	02 Route by risk and skill Use AI for approved routine work, partners for defined human queues, and specialists for judgment.	03 Verify and learn Check the outcome, then update guidance, evaluations, or the product backlog.
--	--	--

Who does what

AI and product help Best for approved, repeatable requests. Does guide or act within tested limits, retain context, and offer a human handoff. Feeds back failed answers and missing guidance.	Partner teams Best for repeatable human cases and coverage across time zones. Do solve cases using approved guidance, staff defined queues, and coach against one quality standard. Feed back recurring customer pain and unclear handoffs.	Internal specialists Best for sensitive cases, unclear policy, incidents, and product defects. Do make judgment calls and own final escalations. Own routing, evaluations, partner governance, and Product and Engineering priorities.
--	---	--

How partners evolve

Partners begin with defined queues and matched quality reviews. As the system improves, they add trained backup coverage, contribute examples to the knowledge and evaluation loop, and earn more volume when durable resolution and rework support it. OpenAI keeps policy, AI release rules, and final escalation decisions.

Start small, then expand the paths that hold up.

First 90 days	Join issue and outcome data, set the quality and capacity baseline, and pilot two low-risk journeys with human review.
Year 1	Improve product help and partner routing on proven journeys. Set one quality rubric across partners and train backup coverage.
Year 2	In Year 2, I would take proven paths across more products and channels. Cross-trained partners would address recurring issue clusters and bring patterns into Product and Engineering plans.
Year 3	Run one linked AI and human support system. Give partners more well-defined volume when outcomes support it, while specialists focus on complex cases and system improvement.

AI answer or policy error Stop the affected route, send cases to a human owner, review the cause, and update the evaluation.	Uneven partner quality Compare similar case mixes, calibrate audits across sites, and shift volume only after results hold.	Hidden demand or capacity gaps Track repeat contact, backlog age, qualified hours, attrition, no-shows, and late arrivals.
--	---	--

Action plan

Status colors show illustrative Day 0 readiness for this exercise. They describe a proposed plan before execution begins. Red means a blocking input, yellow means design needed, and green means ready to start.

ACTION	OWNER ROLE	TIMELINE	ILLUSTRATIVE STATUS
Connect issue history, automation results, and repeat contact.	Support Ops + Analytics	Days 0-30	● Red • data access
Agree on issue types, quality rules, and partner scorecards.	Support Ops + Partner Lead	Days 0-60	● Yellow • design
Pilot two low-risk journeys with human handoff.	Product + Engineering + Support	Days 30-90	● Yellow • design
Map trained backup staffing and specialist escalation.	Workforce Planning + Partners	Year 1	● Green • ready
Expand proven journeys and fix repeat issue clusters.	Support + Product + Engineering	Years 2-3	● Yellow • pilot evidence

What I would ask leadership to approve Give Support Operations linked issue data, two pilot journeys, a Product and Engineering partner, and authority to set common partner quality rules. In 90 days, I would return with an observed capacity plan, a decision on each pilot, and the next set of product fixes.

Judge the customer outcome, then each contributor.

The lead measure is durable resolution by issue type. The customer gets a correct answer with no repeat contact about that issue within the appropriate review window. A shared issue ID shows how automation, partners, and internal specialists contributed without counting a transfer as a fresh resolution.

QUESTION	MEASURES TO REVIEW	WHAT WOULD CHANGE MY PLAN
Did the customer get help?	Durable resolution, repeat contact, audited answer quality, severe errors, time to resolution, and customer effort.	Rising recontact or severe errors would stop expansion on that route and trigger case review.
Can capacity keep up?	Arrivals by issue type, qualified hours, backlog age, occupancy, handle time, partner attrition, no-show rate, and late rate.	Persistent aged backlog would shift schedules, trained backup staffing, or specialist capacity.
Does each step add value?	Eligible AI use, verified AI resolution, handoff quality, partner durable resolution, specialist rework, and cost per durable resolution.	Higher AI use with weaker outcomes would send the journey back to design.
Where is repeat demand coming from?	Issue coverage across products and channels, journey starts, successful routing, resolution, and repeat-contact clusters.	Recurring clusters would move into a Product and Engineering fix queue.

How I would adapt my Meta scorecard			
Revenue.		Coverage.	
Gross and incremental revenue help prioritize business customer pain, while durable resolution remains the Support goal.		Breadth becomes product, channel, and issue coverage. Depth becomes qualified staffing for each queue.	
Solution score.		Quality.	
Pitched calls per book and solutions per call become the share of eligible issues with a relevant next step and a verified outcome.		Tailored solutions and commercial conversations become answer relevance, context use, and clear ownership for business customers.	
Workforce.		Journey funnel.	
Attrition, no-shows, and late arrivals still show staffing reliability. Very short contacts prompt a quality sample.		Attempted, connected, pitched, and actioned become issue start, routing, guidance delivered, and confirmed resolution. Segment by entry channel.	
Tooling.			
Ads AI adoption becomes support-tool use among eligible agents, paired with audited quality and customer outcomes.			

Compare like-for-like issue types, risk, channel, and customer groups. Report the end-to-end result alongside AI, partner, and specialist contributions; keep transfers and reopened issues attached to the original issue record.

The planning assumptions behind this recommendation.

Starting point. The exercise supplies 1 billion weekly users, about 150k weekly issues, 60% automated resolution, and 40% human assistance.

Case complexity. Human effort per case rises to 1.10×, 1.20×, and 1.30× because repeatable issues may leave the human queue first. One effort unit equals today's average human case.

Capacity conversion. The model estimates effort before translating it into headcount. Staffing requires handle time, qualified hours, occupancy, shrinkage, case mix, and service windows.

Demand growth. Weekly users rise to 1.15×, 1.32×, and 1.50× across Years 1–3. Issues per user rise to 1.08×, 1.16×, and 1.25×. These illustrative steps are my planning assumptions.

Automation path. Verified resolution rises to 65%, 70%, and 75% only in this planning case. The exercise provides no target. Count an issue as resolved only after an appropriate repeat-contact window and quality check.

Contribution and risk. A linked issue record attributes each step while preserving the final customer outcome. Human owners review severe errors and expansion decisions.

External source [OpenAI's published support model](#). Its discussion of support surfaces, knowledge, evaluations, and frontline feedback informs the operating design.

This is my independent recommendation for a hypothetical exercise. I have no access to OpenAI internal support data.