

Support at 1 billion weekly users

Lots of organizations deal with scale. Fewer deal with scale and hypergrowth.

The exercise starts with 150,000 issues each week, including 60,000 that need a person. Even a modest rise in issues per user creates substantial new demand. My three-year planning case reaches about 281,000 weekly issues. If automation resolves only 75%, the harder human queue needs about 52% more effort. That is the capacity risk I would plan against.

1 billion weekly users	150k weekly issues	60% automated resolution	60k human-assisted issues
--	--	--	---

“Support has never really been about replying to just tickets. It’s about whether people get what they need, whether it actually serves them well.”

Glen Worthington, Head of User Ops

The three-year vision

Support spots needs inside the product, completes approved actions, and turns every exception into better guidance and product fixes. People own sensitive decisions and AI quality.


My recommendation

Plan toward 88% verified automated resolution as an illustrative stretch case. At modeled Year 3 demand, human cases fall to about 34k weekly and weighted human workload falls to about 44k current-case equivalents. Use the hours AI frees for harder cases, quality review, and repeat-issue fixes. Change staffing only after customer outcomes and coverage hold.

The starting figures come from the exercise. The three-year path is my planning scenario, separate from any OpenAI target.

AI absorbs growth; people handle fewer, harder cases.

The bars show where weekly issues go as demand grows. The stretch case assumes enough eligible issues, approved AI actions, and a verified result that lasts.



HUMAN CASES

60k → 34k per week

This counts issues that reach a person. The case count falls 43% in the stretch case.

HUMAN WORKLOAD

60k → 44k equivalents

Each Year 3 case takes 1.3 times today's average effort. Thus $34k \times 1.3 \approx 44k$ equivalents, about 27% below today's workload.

What must be true

At Year 3 demand, automation needs to exceed about 84% to bring weighted human workload below today's level. I would confirm the eligible issue mix in the first 90 days, then expand only when lasting resolution, safety, and backlog remain sound.

A current-case equivalent measures workload for planning. Staffing still requires measured hours, and the appendix shows every assumption and the 75% downside case.

Build the system before changing staffing.

These three investments run together. Each one reduces avoidable demand or makes the next AI and human resolution more reliable.

PRODUCT Prevent repeat issues Use support patterns to fix confusing product paths and place help where customers get stuck. Expected gain Fewer issues enter Support.	AI ACTIONS Complete approved tasks Give AI bounded actions, current guidance, identity checks, and a full-context handoff to people. Expected gain More eligible issues stay resolved.	PEOPLE Turn exceptions into fixes Give partners and specialists time to review failures, improve guidance, and test new paths. Expected gain Hard cases improve the system.
AI and product help Handle approved routine issues and keep an auditable record.	Partner teams Handle defined exceptions, protect coverage, and flag repeat failures.	Internal specialists Own sensitive decisions, policy, AI release rules, and product fixes.
Capacity decision Use the hours freed by AI for harder cases, quality review, and product fixes first. Training and backup coverage add cost early, so I would adjust staffing through attrition or partner contracts only after measured workload, coverage, customer outcomes, and cost support the change.		

Prove each step before expanding AI.

First 90 days

I would map eligible issues and cost, pilot invoice lookup and product how-to guidance with human fallback, and return with a go or stop decision and capacity forecast.

Year 1

Support expands passing paths toward 70% verified automation and moves freed hours to complex cases and review.

Year 2

Product and Engineering extend approved actions toward 80%, fix repeat causes in product, and retrain partners.

Year 3

The team tests the 88% stretch case and adjusts vendor capacity only after sustained quality and workload gains.

What I need from leadership

- Linked issue, repeat-contact, and cost data across channels.
- Named Product, Engineering, and policy owners.
- Training time and flexible partner terms.

How I would judge progress

- Durable resolution and severe errors show customer quality.
- Human workload and backlog show capacity.
- Cost per lasting resolution shows the full economics.

Decision rule

Track one issue across AI, partners, and internal specialists, and count a lasting resolution once. Stop an affected AI action after a severe error or rising repeat contact. Keep human coverage when workload stays high.

The detailed owners, scorecard, pilot gates, and scenario math remain in the appendix for questions.

The math behind the stretch case.

PLANNING CASE	TODAY	YEAR 1	YEAR 2	YEAR 3
Weekly issues	150k	186k	230k	281k
Verified automated resolution	60%	70%	80%	88%
Human cases per week	60k	56k	46k	34k
Effort per human case	1.0×	1.1×	1.2×	1.3×
Weighted human workload	60k	61k	55k	44k

Starting point. The exercise supplies 1 billion weekly users, about 150k weekly issues, 60% automated resolution, and 40% human assistance.

Case complexity. Each remaining human case takes 1.10×, 1.20×, and 1.30× today's average effort across Years 1–3. Weighted workload equals human cases multiplied by case difficulty.


Staffing conversion. A staffing plan needs measured handle time, qualified hours, occupancy, time away from queues, service windows, and a transition plan. The model's equivalent cases are planning units.

Demand growth. User factors rise to 1.15×, 1.32×, and 1.50× across Years 1–3. Issues per user rise to 1.08×, 1.16×, and 1.25×. These are illustrative assumptions.

Break-even. At Year 3 demand, verified automation must exceed about 84% to lower weighted human workload below today's level. At 75%, people handle about 70k weekly cases and 91k equivalent cases of work.

Eligibility gate. I would measure eligible issue volume by type before endorsing the 88% target. The model gives no credit for product fixes that prevent an issue from arising.

How the model runs and recovers.



My approach Diagnose demand and case complexity, design the operating model, and deliver verified gains in stages.

Read the customer result and each contributor.

OPERATING LENS	MEASURE	LEADERSHIP RESPONSE
Quality	Durable resolution, repeat contact, audited answer quality, severe errors.	Stop the affected AI path if severe errors or repeat contact rise.
Productivity	Verified automation, human effort per 1,000 issues, time to resolution.	Expand only when workload falls with customer quality intact.
Capacity	Qualified hours, backlog age, coverage by issue type, attrition, no-shows.	Move hours to higher-skill tasks first and protect service coverage.
Cost to serve	Cost per durable resolution, including handoffs, rework, and transition costs.	Change vendor volume only when full cost improves.

One issue history shows each contribution		
AI	Partners	Internal teams
Eligible coverage, verified resolution, handoff quality, and error rate.	Lasting resolution, audit quality, staffing reliability, and rework on comparable cases.	Complex-case outcomes, escalation quality, and repeat demand removed through fixes.

Proposed pilots. Invoice lookup and product how-to guidance are examples. A Day 30 review would confirm eligible volume, risk, baseline, and the right review window.

Illustrative Day 90 gates. I would look for a 5-point gain in durable resolution, 10% lower human effort per eligible issue, and zero severe errors before expansion. These are planning gates, separate from OpenAI commitments.

External source [OpenAI's published support model](#). Its support surfaces, knowledge, evaluations, and frontline feedback inform this proposal.

This is my independent recommendation for a hypothetical exercise. I have no access to OpenAI internal support data.