

Support at 1 billion weekly users

Lots of organizations deal with scale. Fewer deal with scale and hypergrowth.

The exercise starts with 150,000 support issues each week. People handle 60,000 of them.

In my three-year planning scenario, weekly issues reach about 281,000. AI would need to resolve far more issues, while people handle the harder cases that remain.

1 billion weekly users	150k weekly issues	60% automated resolution	60k human-assisted issues
--	--	--	---

“Support has never really been about replying to just tickets. It’s about whether people get what they need, whether it actually serves them well.”

Glen Worthington, Head of User Ops

The three-year vision

Support spots needs inside the product, completes approved actions, and turns every exception into better guidance and product fixes. People own sensitive decisions and AI quality.


My recommendation

Plan toward 88% verified automated resolution as an illustrative stretch case. At modeled Year 3 demand, issues reaching people fall from 60k to about 34k each week. Move people toward harder cases, quality review, and repeat-issue fixes. Change staffing only after customer outcomes and coverage hold.

The starting figures come from the exercise. The three-year path is my planning scenario, separate from any OpenAI target.

AI absorbs growth and people handle fewer, harder cases.

The bars show where weekly issues go as demand grows. The stretch case assumes enough eligible issues, approved AI actions, and a verified result that lasts.



CASES REACHING PEOPLE

60k → 34k per week

The case count falls 43% in the stretch case, even as total issues rise.

CAPABILITY NEEDED

Stronger specialist skills

As AI resolves more routine issues, people focus on sensitive cases, exceptions, and quality review.

What must be true

Enough issues must qualify for safe AI resolution, and customers must stay helped after the interaction ends. I would check the eligible issue mix in the first 90 days, then expand only when lasting resolution, safety, and backlog remain sound.

The appendix shows the demand assumptions and the case count if automation reaches only 75% in Year 3.

Build the system before changing staffing.

These three investments run together. Each one reduces avoidable demand or makes the next AI and human resolution more reliable.

PRODUCT Prevent repeat issues Use support patterns to fix confusing product paths and place help where customers get stuck. Expected gain Fewer issues enter Support.	AI ACTIONS Complete approved tasks Give AI bounded actions, current guidance, identity checks, and a full-context handoff to people. Expected gain More eligible issues stay resolved.	PEOPLE Turn exceptions into fixes Give partners and specialists time to review failures, improve guidance, and test new paths. Expected gain Hard cases improve the system.
AI and product help Handle approved routine issues and keep an auditable record.	Partner teams Handle defined exceptions, protect coverage, and flag repeat failures.	Internal specialists Own sensitive decisions, policy, AI release rules, and product fixes.
Capacity decision As AI handles more eligible issues, I would move people toward harder cases, quality review, and product fixes first. Training and backup coverage add cost early, so I would adjust staffing through attrition or partner contracts only after case volume, coverage, customer outcomes, and cost support the change.		

Prove each step before expanding AI.

First 90 days

Map eligible issues, record quality and cost baselines, and pilot invoice lookup and product how-to guidance. Decide on each pilot using lasting resolution, repeat contact, severe errors, and cases reaching people.

Year 1

Target 70% verified automation and about 56k weekly cases reaching people. Keep customer outcomes and backlog at least as strong as the baseline before expanding.

Year 2

Target 80% verified automation and about 46k weekly cases reaching people. Show lower cost per lasting resolution while repeat contact and backlog stay at or below baseline.

Year 3

Test 88% verified automation and about 34k weekly cases reaching people. Adjust vendor capacity only after two consecutive reviews confirm strong quality, reliable coverage, and lower cost.

Check every year Report verified AI resolution, cases reaching people, lasting resolution, repeat contact, severe errors, backlog age, and cost per lasting resolution by issue type. The numeric targets above are my planning assumptions, separate from OpenAI commitments.

What I need from leadership

- Linked issue, repeat-contact, and cost data across channels.
- Named Product, Engineering, and policy owners.
- Training time and flexible partner terms.

How I would judge progress

- Durable resolution and severe errors show customer quality.
- Human case volume and backlog show capacity.
- Cost per lasting resolution shows the full economics.

Decision rule

Track one issue across AI, partners, and internal specialists, and count a lasting resolution once. Stop an affected AI action after a severe error or rising repeat contact. Keep human coverage if queues stay high.

The detailed owners, scorecard, pilot gates, and scenario math remain in the appendix for questions.

The math behind the stretch case.

PLANNING CASE	TODAY	YEAR 1	YEAR 2	YEAR 3
Weekly issues	150k	186k	230k	281k
Verified automated resolution	60%	70%	80%	88%
Human cases per week	60k	56k	46k	34k

Starting point. The exercise supplies 1 billion weekly users, about 150k weekly issues, 60% automated resolution, and 40% human assistance.

Case complexity. As AI resolves more routine issues, I expect the remaining human queue to include more sensitive cases and exceptions. Partners and specialists would need training, current guidance, and clear escalation paths.


Staffing decision. I would use observed case volume, coverage, backlog, customer outcomes, and cost before adjusting staffing. The forecast assigns no headcount target.

Demand growth. User factors rise to 1.15×, 1.32×, and 1.50× across Years 1–3. Issues per user rise to 1.08×, 1.16×, and 1.25×. These are illustrative assumptions.

Slower automation. If verified automation reaches only 75% in Year 3, about 70k issues would reach people each week, above today's 60k. I would keep coverage in place and investigate the issue mix before adjusting staffing.

Eligibility gate. I would measure eligible issue volume by type before endorsing the 88% target. The model gives no credit for product fixes that prevent an issue from arising.

How the model runs and recovers.



My approach Diagnose demand and case complexity, design the operating model, and deliver verified gains in stages.

Read the customer result and each contributor.

OPERATING LENS	MEASURE	LEADERSHIP RESPONSE
Quality	Durable resolution, repeat contact, audited answer quality, severe errors.	Stop the affected AI path if severe errors or repeat contact rise.
Productivity	Verified automation, human case volume, time to resolution.	Expand AI only when customer quality holds and fewer cases reach people.
Capacity	Qualified hours, backlog age, coverage by issue type, attrition, no-shows.	Move people to higher-skill tasks first and protect service coverage.
Cost to serve	Cost per durable resolution, including handoffs, rework, and transition costs.	Change vendor volume only when full cost improves.

One issue history shows each contribution		
AI	Partners	Internal teams
Eligible coverage, verified resolution, handoff quality, and error rate.	Lasting resolution, audit quality, staffing reliability, and rework on comparable cases.	Complex-case outcomes, escalation quality, and repeat demand removed through fixes.

Proposed pilots. Invoice lookup and product how-to guidance are examples. A Day 30 review would confirm eligible volume, risk, baseline, and the right review window.

Illustrative Day 90 gates. I would look for a 5-point gain in durable resolution, 10% fewer eligible cases reaching people, and zero severe errors before expansion. These are planning gates, separate from OpenAI commitments.

External source [OpenAI's published support model](#). Its support surfaces, knowledge, evaluations, and frontline feedback inform this proposal.

This is my independent recommendation for a hypothetical exercise. I have no access to OpenAI internal support data.