

Support at 1 billion weekly users

Lots of organizations deal with scale. Fewer deal with scale and hypergrowth.

The exercise already starts with 150,000 issues and 60,000 human-assisted cases each week. A small increase in issue frequency creates a large amount of additional demand.

In my three-year planning case, issues approach 281,000 each week. If verified automation reaches only 75%, the harder human queue requires about 52% more effort, raising the risk of slower answers and repeat contacts. I would aim for a higher verified resolution rate while keeping people in control of sensitive decisions.

1 billion weekly users	150k weekly issues	60% automated resolution	60k human-assisted issues
--	--	--	---

“Support has never really been about replying to just tickets. It’s about whether people get what they need, whether it actually serves them well.”

Glen Worthington, Head of User Ops

The vision

What this looks like in 3 years

Support spots common needs and resolves approved problems inside the product before they become tickets, while people govern AI, handle sensitive cases, and remove the causes of repeat issues.

My recommendation

Use an illustrative path from 60% to 88% verified automated resolution by Year 3. At modeled demand, human cases fall from 60k to about 34k weekly and human effort falls from 60k to 44k units; I would redeploy released capacity first.

How I approached this assignment

Diagnose Demand and case complexity	Design AI, people, and control points	Deliver Verified gains in stages
---	---	--


The starting figures come from the hypothetical exercise. The three-year path is my planning scenario, separate from any OpenAI target.

AI takes more cases as human cases fall.

My stretch case assumes weekly issues rise from 150k to 281k and the average human case becomes 1.3× harder. Verified automated resolution rises in stages to 88%, which is enough to lower both human case volume and human effort.

Where the weekly issues go

Illustrative weekly volume. Bar length shows total demand; green is resolved through automation and gold goes to a person.



What the chart shows

Automation handles about 157k more weekly issues by Year 3, while human cases fall by about 26k. The chart shows case flow; the staffing decision depends on the effort and quality checks below.

About 84% is the effort break-even point

At Year 3 demand and 1.3× effort per human case, verified automation must exceed about 84% to bring human effort below today's level.

88% is the stretch case

Human cases reach about 34k weekly and effort reaches about 44k units, 27% below today. At 75% automation, effort would rise to about 91k units.

Capacity must change with proof

Keep qualified coverage while AI takes on more work. Redeploy released hours into review, difficult cases, and product fixes; reset staffing through attrition or partner contracts only after durable resolution, backlog, and safety hold.

Capability must change first

Give AI approved actions, current knowledge, tests, traces, and a fast human handoff. Train partners to handle exceptions and surface failures, while internal specialists own policy, evaluations, and high-risk decisions.

Scenario logic Year 3 uses 1.5× weekly users, 1.25× issues per user, and 1.3× effort per human case. The path gives no credit for preventing issues; product fixes could improve demand further.

Eligibility is the gate. The 88% path requires enough issues to qualify for safe automated resolution; I would measure eligible volume by issue type before endorsing this target.

Three moves make the 88% path credible.

[OpenAI's published support model](#) describes connected support surfaces, living knowledge, evaluations, and frontline feedback. I would extend that system so it resolves more needs inside the product and uses verified outcomes to reduce repeat demand.

PRODUCT

Prevent repeat issues

Use support patterns to fix confusing flows and place help where customers get stuck.

Reward

Fewer issues reach Support as usage grows.

Risk to manage

Confirm the product change actually lowers repeat contact.

AI ACTIONS

Complete approved work

Move beyond answers to bounded actions such as invoice lookup, with identity checks, audit trails, and human handoff.

Reward

More issues stay resolved without a human ticket.

Risk to manage

A wrong action needs a stop rule and recovery path.

PEOPLE

Make frontline teams builders

Give partners and specialists time to flag failures, improve guidance, and create test cases.

Reward

Each exception improves the next automated and human answer.

Risk to manage


Protect coverage while people learn these new tasks.

Investment order

Start with shared issue history, evaluations, and safe action tools. Expand AI coverage and change partner capacity only after the customer outcome holds for comparable cases.

AI resolves the routine; people govern the system.

One issue history follows the customer across product help, partners, and specialists. Support gives AI approved scope, and people keep clear approval, escalation, and rollback rights.



Prove each step before changing capacity.

Readiness colors are illustrative Day 0 planning statuses. Red needs a blocking input, yellow needs design or pilot evidence, and green is ready to start.

WHEN	ACTION AND EVIDENCE	OWNER	ILLUSTRATIVE STATUS
Week 1	Name the pilot owners and agree the issue types to baseline.	Support Ops	● Green · ready to start
First 90 days	Link issue outcomes, baseline case effort, and pilot invoice lookup plus product how-to guidance.	Support Ops + Product + Engineering	● Red · data access
Year 1	Expand passing journeys toward 70% verified automation; retrain and redeploy released hours.	Support Ops + Partners	● Yellow · pilot evidence
Year 2	Extend approved actions across more surfaces toward 80%; correct repeat issue patterns in product.	Product + Engineering + Support	● Yellow · design
Year 3	Test the 88% stretch case and adjust contracted capacity only after sustained quality and workload gains.	Support Ops + Partner Lead + Finance	● Red · proof needed

What I need from leadership • Linked issue, cost, and repeat-contact data across channels. • Named Product, Engineering, and policy owners for approved AI actions. • Training capacity and flexible partner terms for the transition.	What I will bring by Day 90 • A baseline and a decision on two bounded pilots. • Illustrative gates of a 5-point gain in durable resolution, 10% less human effort per eligible issue, and zero severe errors. • A capacity forecast and the first product fixes.
---	--

Severe error Stop that action and route affected cases to people.	More repeat contact Freeze expansion, repair guidance or tools, and rerun tests.	Human effort stays high Keep staffing in place and revisit case mix and handoffs.
---	--	---

Scale AI only when customer outcomes hold.

The lead outcome is durable resolution by issue type. The customer gets the right result without returning with that issue during an appropriate review window. One issue ID keeps AI, partner, and specialist contributions visible across handoffs.

OPERATING LENS	MEASURE	LEADERSHIP RESPONSE
Quality	Durable resolution, repeat contact, audited answer quality, severe errors.	Pause the affected AI path if recontact or severe errors rise.
Productivity	Verified automated resolution, human effort per 1,000 issues, time to resolution.	Expand only when human effort falls with quality intact.
Capacity	Qualified hours, backlog age, coverage by issue type, attrition, no-shows.	Redeploy hours first; change staffing after coverage remains sound.
Cost to serve	Cost per durable resolution, including handoffs, rework, and transition costs.	Change vendor volume only when the end-to-end cost improves.

Read the total result and each contributor		
AI	Partners	Internal teams
Eligible coverage, verified resolution, handoff quality, and error rate.	Durable resolution, audit quality, staffing reliability, and rework on comparable cases.	Complex-case outcomes, escalation quality, and repeat demand removed through fixes.

The 88% path earns expansion only when durable resolution, safety, backlog, and cost improve together. Compare similar issue types and customer groups before changing capacity.

The numbers behind the stretch case.

Starting point. The exercise supplies 1 billion weekly users, about 150k weekly issues, 60% automated resolution, and 40% human assistance.

Case complexity. Human effort per case rises to 1.10×, 1.20×, and 1.30× as repeatable issues leave the queue. One unit equals today's average human case.

Effort threshold. At Year 3 volume, about 84% verified automation holds human effort below today's level. At 75%, human cases reach about 70k and effort about 91k units.

Staffing conversion. Effort units guide planning. Staffing requires observed handle time, qualified hours, occupancy, shrinkage, service windows, and a transition plan.

External source [OpenAI's published support model](#). Its support surfaces, knowledge, evaluations, and frontline feedback inform this proposal.

This is my independent recommendation for a hypothetical exercise. I have no access to OpenAI internal support data.

Demand growth. User factors rise to 1.15×, 1.32×, and 1.50× across Years 1–3. Issues per user rise to 1.08×, 1.16×, and 1.25×. These are my illustrative assumptions.

AI path. Verified automated resolution rises to 70%, 80%, and 88% in this stretch case. The exercise provides no OpenAI target. Count a resolution only after quality and repeat-contact checks.

Year-by-year outlook. Today shows 150k issues, 60k human cases, and 60k effort units; Year 1 shows 186k, 56k, and 61k; Year 2 shows 230k, 46k, and 55k; Year 3 shows 281k, 34k, and 44k.

Pilot gates. Invoice lookup and product how-to guidance are proposed examples. The Day 90 gates depend on a Day 30 baseline, risk review, and an appropriate resolution window.