

Support at 1 billion weekly users

Lots of organizations deal with scale. Fewer deal with scale and hypergrowth.

The exercise already starts with 150,000 issues and 60,000 human-assisted cases each week. A small increase in issue frequency creates a large amount of additional demand. In my three-year planning case, issues approach 281,000 each week, while a harder human queue requires 52% more effort even with stronger automation. Without a change in the operating model, slower answers and repeat contacts become a bigger risk.

1 billion weekly users	150k weekly issues	60% automated resolution	60k human-assisted issues
--	--	--	---

“Support has never really been about replying to just tickets. It’s about whether people get what they need, whether it actually serves them well.”

Glen Worthington, Head of User Ops

The vision

What this looks like in 3 years

AI resolves approved routine issues, helps people handle complex cases without losing context, and turns customer interactions into stronger guidance and product fixes, with human judgment for sensitive decisions.


How I approached this assignment

Diagnose Demand and case complexity	Design The operating model	Deliver In measured stages
---	--	--

The starting figures come from the hypothetical take-home exercise. The three-year figures are my planning assumptions.

Weekly demand nearly doubles and human effort still rises.

I use a three-year planning case to show the capacity implications. By Year 3, weekly users reach 1.5× today's exercise baseline, issues per user reach 1.25×, and a human case takes 1.3× today's average effort. Verified automated resolution rises from 60% to 75% in stages.



PLANNING VIEW	WEEKLY ISSUES	VERIFIED AUTOMATION	HUMAN CASES	HUMAN EFFORT UNITS
Today 1.00× users · 1.00× issues/user · 1.00× effort	150k	60%	60k	60k
Year 1 1.15× users · 1.08× issues/user · 1.10× effort	186k	65%	65k	72k
Year 2 1.32× users · 1.16× issues/user · 1.20× effort	230k	70%	69k	83k
Year 3 1.50× users · 1.25× issues/user · 1.30× effort	281k	75%	70k	91k

Why these assumptions A larger user base raises demand. More products and business use cases can lift issue frequency. Automation can remove easier cases first, leaving a harder human queue. The multipliers illustrate those pressures; they are my judgment, separate from any OpenAI forecast.

One effort unit equals today's average human case. The figures are rounded estimates. If automation stays at 60% in Year 3, the model reaches about 113k human cases and 146k effort units.

Capacity

Human cases rise from 60k to about 70k each week, while effort rises from 60k to 91k units. I would plan qualified hours and flexible coverage by case type before converting effort into staffing.

Capability

The remaining cases require stronger judgment and deeper product knowledge. I would train partners for defined queues and give internal specialists the tools and time to handle sensitive issues and fix recurring causes.

AI-led Support that learns from every interaction.

[OpenAI's published support model](#) connects customer surfaces, living knowledge, evaluations, and frontline feedback. I would extend that approach across the next three years. AI should resolve approved routine issues, help people handle complex ones, and turn repeated failures into better guidance and product changes.

The strategy

Where to play and how to win

Where to play

- Start with common, low-risk issues where an approved answer or action can be checked.
- Keep sensitive, unclear, and high-risk issues with trained people.

How to win

- Keep one issue history across AI, partners, and specialists.
- Give each team current guidance and test answers against customer outcomes.
- Send repeat failures to Product and Engineering for lasting fixes.

These three investments run together. Shared issue data and quality checks determine when AI coverage can expand.

CUSTOMER HELP

Resolve routine issues in product

Invest in guided help and approved AI actions with a clear path to a person.

Expected benefit

Faster answers and less repeatable demand.

Near-term cost

A narrow pilot delays broader coverage while quality is checked.

LEARNING LOOP

Improve each answer

Join issue history, living knowledge, evaluations, and frontline feedback.

Expected benefit

Better quality across channels and fewer recurring failures.

Near-term cost

Data and quality setup require time before broad rollout.

HUMAN CAPABILITY

Build for complex cases

Train partners for defined queues and internal specialists for judgment and product defects.

Expected benefit

Reliable coverage as the remaining case mix gets harder.


Near-term cost

Training and backup coverage add expense before productivity improves.

A higher automation rate helps only when customers stay resolved and severe errors remain controlled.

One customer issue, one history, clear owners.

The operating model begins with one issue record. It holds the customer's question, previous attempts, risk level, handoffs, and final outcome. Routing then sends the issue to the right capability instead of forcing every request through one queue.



Who does what

AI and product help Best for approved, repeatable requests. Does guide or act within tested limits, retain context, and offer a human handoff. Feeds back failed answers and missing guidance.	Partner teams Best for repeatable human cases and coverage across time zones. Do solve cases using approved guidance, staff defined queues, and coach against one quality standard. Feed back recurring customer pain and unclear handoffs.	Internal specialists Best for sensitive cases, unclear policy, incidents, and product defects. Do make judgment calls and own final escalations. Own routing, evaluations, partner governance, and Product and Engineering priorities.
--	---	--

How partners evolve

Partners begin with defined queues and matched quality reviews. As the system improves, they add trained backup coverage, contribute examples to the knowledge and evaluation loop, and earn more volume when durable resolution and rework support it. OpenAI keeps policy, AI release rules, and final escalation decisions.

Start small, then expand the paths that hold up.

First 90 days	- Join issue and outcome data, then set the quality and capacity baseline. - Pilot two low-risk journeys with human review and clear stop rules.
Year 1	- Improve product help and partner routing on journeys that pass the pilot. - Use one quality rubric across partners and train backup coverage.
Year 2	- Expand proven journeys across more products and channels. - Cross-train partners and send recurring issue patterns to Product and Engineering.
Year 3	- Run one linked AI and human Support system. - Give partners more defined volume when outcomes support it, while specialists handle complex cases and system fixes.

Largest risks and responses

These are the conditions that could slow or stop expansion; each one has a response.

AI answer or policy error Stop the affected route, send cases to a human owner, review the cause, and update the evaluation.	Uneven partner quality Compare similar case mixes, calibrate audits across sites, and shift volume only after results hold.	Hidden demand or capacity gaps Track repeat contact, backlog age, qualified hours, attrition, no-shows, and late arrivals.
--	---	--

Make the first 90 days accountable.

Status colors show illustrative Day 0 readiness for this exercise. Red means a blocking input, yellow means design needed, and green means ready to start.

ACTION	OWNER ROLE	TIMELINE	ILLUSTRATIVE STATUS
Connect issue history, automation results, and repeat contact.	Support Ops + Analytics	Days 0-30	● Red • data access
Agree on issue types, quality rules, and partner scorecards.	Support Ops + Partner Lead	Days 0-60	● Yellow • design
Pilot invoice or receipt lookup and product how-to guidance with human handoff.	Product + Engineering + Support	Days 30-90	● Yellow • design
Map trained backup staffing and specialist escalation.	Workforce Planning + Partners	Year 1	● Green • ready
Expand proven journeys and fix repeat issue clusters.	Support + Product + Engineering	Years 2-3	● Yellow • pilot evidence

What I would ask leadership to approve

The ask

- Give Support Operations linked issue and repeat-contact data.
- Name Product and Engineering owners for two pilot journeys covering invoice or receipt lookup and a common product how-to question.
- Allow one quality standard and clear partner routing rules.

The journeys and targets are illustrative until volume, risk, and Day 30 baselines are reviewed.

Return by Day 90

- Move from no shared journey baseline to measured results by Day 30.
- Target a 5-point gain in 7-day durable resolution and 10% less human effort per eligible issue versus that baseline, with zero severe errors.
- Bring a capacity plan, a decision on each pilot, and the top product fixes.

Use four operating lenses, anchored in customer resolution.

The lead outcome is durable resolution by issue type. The customer gets a correct answer and has no repeat contact about that issue during an appropriate review window. I would review quality, productivity, capacity, and cost together so that improvement in one area cannot hide a decline in another.

OPERATING LENS	MEASURES TO REVIEW	WHAT WOULD CHANGE MY PLAN
Quality	Durable resolution, repeat contact, audited answer relevance, handoff quality, and severe errors.	If use rises while resolution falls, pause expansion and review the failed cases.
Productivity	Eligible issues completed per qualified hour, time to resolution, successful journey completion, and tool use among eligible agents.	If tool use stays low or handoffs add time, change the guidance and workflow before expanding.
Capacity	Arrivals and case mix, backlog age, qualified hours, coverage across issues and channels, attrition, no-shows, and late arrivals.	If aged backlog rises, shift schedules, train backup coverage, or add specialist capacity.
Cost to serve	Cost per durable resolution by comparable issue type, including handoffs, rework, and backup coverage.	If lower unit cost comes with more rework or errors, change the route or volume allocation.

How I would read each contributor		
Automation	Partners	Internal teams
Track eligible coverage, verified resolution, handoff quality, and severe errors.	Compare durable resolution, audited quality, staffing reliability, and rework for similar cases.	Track complex-case resolution, escalation quality, product fixes, and repeat demand removed.

Use one issue ID across every handoff and compare like-for-like case mixes. For business customers, affected accounts and revenue at risk help prioritize fixes; durable resolution remains the Support outcome.

The planning assumptions behind this recommendation.

Starting point. The exercise supplies 1 billion weekly users, about 150k weekly issues, 60% automated resolution, and 40% human assistance.

Case complexity. Human effort per case rises to 1.10×, 1.20×, and 1.30× because repeatable issues may leave the human queue first. One effort unit equals today's average human case.

Capacity conversion. The model estimates effort before translating it into headcount. Staffing requires handle time, qualified hours, occupancy, shrinkage, case mix, and service windows.

Pilot gates. Invoice or receipt lookup and product how-to guidance are proposed examples. The Day 90 targets are illustrative and depend on a Day 30 baseline and risk review.

Demand growth. Weekly users rise to 1.15×, 1.32×, and 1.50× across Years 1–3. Issues per user rise to 1.08×, 1.16×, and 1.25×. These illustrative steps are my planning assumptions.

Automation path. Verified resolution rises to 65%, 70%, and 75% only in this planning case. The exercise provides no target. Count an issue as resolved only after an appropriate repeat-contact window and quality check.

Contribution and risk. A linked issue record attributes each step while preserving the final customer outcome. Human owners review severe errors and expansion decisions.

External source [OpenAI's published support model](#). Its discussion of support surfaces, knowledge, evaluations, and frontline feedback informs the operating design.

This is my independent recommendation for a hypothetical exercise. I have no access to OpenAI internal support data.