

Support at 1 billion weekly users

Lots of organizations deal with scale. Fewer deal with scale and hypergrowth.

The exercise starts with 150,000 support issues each week. People handle 60,000 of them.

My three-year planning scenario reaches 500,000 weekly issues while AI cuts the share handled by people from 40% today to 10% in Year 3.

1 billion weekly users	150k weekly issues	60% automated resolution	60k human-assisted issues
--	--	--	---

THREE-YEAR STRETCH CASE

~\$118M

Potential net vendor cost saved against a forecast that keeps AI resolution at 60%. This high case depends on verified AI results, case handling time, and vendor contracts.

[See the low, medium, and high cases.](#)

“Support has never really been about replying to just tickets. It’s about whether people get what they need, whether it actually serves them well.”

Glen Worthington, Head of User Ops

The three-year vision

AI transforms Support from a queue that responds to issues into a connected system that anticipates needs, resolves approved problems inside the product, and learns from every exception.


My recommendation

Plan toward 90% verified AI resolution by Year 3 as a stretch target, which would reduce the human share of weekly issues from 40% to 10%. Move trained people into complex cases, AI review, and repeat-issue fixes first. Reduce planned vendor seats only after customers stay helped and service coverage stays strong.

The starting figures come from the exercise. The three-year path is my planning scenario, separate from any OpenAI target.

AI absorbs hypergrowth and people handle fewer cases.

The bars show where weekly issues go as demand grows. The stretch case assumes enough eligible issues, approved AI actions, and a verified result that lasts.



CASES REACHING PEOPLE

?

40% → 10%

The human share falls across the plan. Weekly human cases rise in Year 1, then fall to 50k by Year 3.

CAPABILITY NEEDED

?

Interoperability across teams

Sales and Product Engagement bring customer context, while Support Delivery, Operations, Product, and Trust & Safety coordinate resolution and AI review.

OpenAI's [published support model](#) informs this vision. The growth and AI targets are my exercise assumptions.

Change the product, people, and process together.

PRODUCT

Resolve more inside the product

Product and Product Engagement use Support signals to remove repeat issues. AI gives approved help and actions at the point of need.

PEOPLE

Give people the harder work

Support Delivery trains specialists and partners for exceptions. Trust & Safety reviews sensitive paths, while AI brings context and suggests next steps.

PROCESS

Run one learning loop

Operations links case history, quality checks, and vendor plans. AI results feed Product fixes, coaching, and quarterly decisions.

First 90 days

• Map eligible issues and baseline quality

• Pilot invoice lookup and product guidance

• Confirm owners with Product and Trust & Safety

Year 1

• 70% verified AI resolution

• 30% human case share

• Protect coverage during pilot expansion

Year 2

• 82% verified AI resolution

• 18% human case share

• Extend proven paths across products

Year 3

• 90% verified AI resolution

• 10% human case share

• Adjust vendor plans after quality checks

Check every quarter

Quality

(Durable resolution , repeat contact, severe AI errors, tailored guidance)

Productivity

(Verified AI resolution, approved AI tool use, time to resolution)

Capacity

(Coverage breadth and depth, backlog age, attrition, no-shows)

Cost to Serve

(Cost per durable resolution, vendor spend, AI run cost, rework)

A \$118M high case, with clear funding gates.

All three cases assume 220k, 330k, and 500k weekly issues across Years 1-3. They compare vendor costs with a forecast that keeps AI resolution at 60%. The figures include proposed AI, tooling, training, and quality-check costs. They exclude any estimate for revenue growth.

LOW · SLOWER AI, RIGID CONTRACTS -\$13.1M AI reaches 82% in Year 3, while only 10% of potential vendor savings reaches the budget. Full \$20M spending would exceed savings.	MEDIUM · AI TARGETS, SLOWER SAVINGS +\$42.4M AI reaches 90% in Year 3. Vendors would have handled 70% of shifted cases, and contract savings arrive more slowly.	HIGH · AI TARGETS, FLEXIBLE CONTRACTS +\$117.5M AI reaches 90% in Year 3. Vendors would have handled 90% of shifted cases, and contracts allow most planned seats to be removed.
---	---	---

The model already uses 30% onshore and 70% offshore vendor costs . That mix is included in every case. [See the inputs and public price references.](#)

Separate location option Moving 100 suitable vendor seats from onshore to offshore could save about \$7.0M per year before transition costs. That is separate from AI savings and carries potential quality risk in language, specialist accuracy, and escalation.	
What I need from leadership Issue and cost data. Confirm which cases AI can resolve and which vendor costs can change. Named partners. Product, Engineering, Trust & Safety, Support Delivery, Operations, and Product Engagement can clear pilot decisions quickly. A capped first-stage budget. Fund the pilots inside the proposed \$4M Year 1 budget, then release more funding after results and Finance review.	How I would judge progress Day 90. Target 10% fewer eligible cases reaching people, a 5-point gain in durable resolution, and zero severe AI errors. Each year. Check AI share against 70%, 82%, and 90%, alongside human share of 30%, 18%, and 10%. Each quarter. Require repeat contact and backlog age at or below baseline, plus vendor savings confirmed by Finance.
REWARD If quality holds and contracts adjust, the high case saves about \$118M net across three years.	RISK AND RESPONSE Wrong AI answers or rigid contracts can erase savings. Stop the affected AI path after a severe error, keep human coverage if queues rise, and release funding in stages.

How the three-year path was built.

PLANNING CASE	TODAY	YEAR 1	YEAR 2	YEAR 3
Weekly issues	150k	220k	330k	500k
Weekly user factor ?	1.00×	1.21×	1.48×	1.83×
Issues per user factor ?	1.00×	1.21×	1.48×	1.83×
Verified AI resolution	60%	70%	82%	90%
Human case share	40%	30%	18%	10%

Starting point. The exercise supplies 1 billion weekly users, about 150k weekly issues, and a 60% automated resolution rate.

Human case mix. As AI resolves more common issues, the remaining cases may need more judgment, current guidance, and stronger escalation paths.


AI-ready issues. Before expanding AI, I would confirm the volume of issues it can handle safely with approved answers or actions, then check that customers stay helped.

Demand growth. I use matching growth factors for users and issues per user. The underlying factors are square roots of each year's issue total divided by 150k, which keeps the displayed volumes consistent after rounding.

Slower AI progress. At 82% AI resolution in Year 3, about 90k weekly cases still reach people. I would keep coverage until the issue mix and service levels support a change.

Staffing. This scenario sets case-share targets and leaves headcount open. Vendor plans follow observed demand, quality, coverage, and contract terms.

One issue history from start to finish.



Inputs behind the financial range.

INPUT	HIGH CASE	REASON
Vendor rep cost ?	\$93.6k onshore \$23.4k offshore	Public prices of \$45 hourly onshore and \$1,950 monthly offshore.
Vendor location mix	30% onshore 70% offshore	Already included in the \$44.46k yearly blended price in every case.
Cases vendors would handle ?	90%	AI takes common cases that partners may otherwise handle. Actual issue ownership must confirm this.
Cases per vendor rep	3,000 yearly	About 12 cases per scheduled day. Case logs must confirm handling time.
Savings that reach the budget ?	40%, 75%, 90%	Contract changes and fewer planned seats take effect over Years 1-3.
AI and rollout spending ?	\$4M, \$6M, \$10M	Illustrative yearly budget. Finance would replace this with a priced plan.

The high case saves \$137.5M in vendor costs against the flat-60% AI forecast and spends \$20M on AI and rollout, leaving \$117.5M net. Today's vendor bill could still rise as total issue volume grows.

Public vendor price references [SupportOps Global](#), [ConnectPro](#), and [PITON-Global](#). OpenAI contract prices may differ from these planning references.