

Support at 1 billion weekly users

Lots of organizations deal with scale. Fewer deal with scale and hypergrowth.

The exercise starts with 150,000 support issues each week. People handle 60,000 of them.

In my three-year planning scenario, weekly issues reach 500,000. AI would need to absorb that growth while people handle fewer, harder cases.

1 billion weekly users	150k weekly issues	60% automated resolution	60k human-assisted issues
--	--	--	---

THREE-YEAR FINANCIAL OPPORTUNITY

~\$118M

Illustrative net vendor cost avoided against a flat 60% automation forecast. It requires 90% verified AI resolution by Year 3, a more involved vendor case mix, and flexible contracts. [See the plan and conditions.](#)

“Support has never really been about replying to just tickets. It’s about whether people get what they need, whether it actually serves them well.”

Glen Worthington, Head of User Ops

The three-year vision

Support spots needs inside the product, completes approved actions, and turns every exception into better guidance and product fixes. People own sensitive decisions and AI quality.


My recommendation

Plan toward 90% verified AI resolution by Year 3 as a stretch target. Even if weekly issues rise to 500k, cases reaching people fall from 60k to about 50k. Move people toward harder cases, quality review, and repeat-issue fixes. Change staffing only after customer outcomes and coverage hold.

The starting figures come from the exercise. The three-year path is my planning scenario, separate from any OpenAI target.

AI absorbs hypergrowth and people handle fewer cases.

The bars show where weekly issues go as demand grows. The stretch case assumes enough eligible issues, approved AI actions, and a verified result that lasts.



CASES REACHING PEOPLE

60k → 50k per week

The queue rises to 66k in Year 1, then falls below today's level by Year 3.

CAPABILITY NEEDED

Stronger specialist skills

As AI resolves more issues, people focus on sensitive cases, exceptions, and quality review.

What must be true

Enough issues must qualify for safe AI resolution, and customers must stay helped after the interaction ends. I would check the eligible issue mix in the first 90 days, then expand only when lasting resolution, safety, and backlog remain sound.

The appendix shows the demand assumptions and the case count if automation reaches only 82% in Year 3.

Build the system before changing staffing.

These three investments run together. Each one reduces avoidable demand or makes the next AI and human resolution more reliable.

PRODUCT Prevent repeat issues Use support patterns to fix confusing product paths and place help where customers get stuck. Expected gain Fewer issues enter Support.	AI ACTIONS Complete approved tasks Give AI bounded actions, current guidance, identity checks, and a full-context handoff to people. Expected gain More eligible issues stay resolved.	PEOPLE Turn exceptions into fixes Give partners and specialists time to review failures, improve guidance, and test new paths. Expected gain Hard cases improve the system.
AI and product help Handle approved routine issues and keep an auditable record.	Partner teams Handle defined exceptions, protect coverage, and flag repeat failures.	Internal specialists Own sensitive decisions, policy, AI release rules, and product fixes.
Capacity decision As AI handles more eligible issues, I would move people toward harder cases, quality review, and product fixes first. Training and backup coverage add cost early, so I would adjust staffing through attrition or partner contracts only after case volume, coverage, customer outcomes, and cost support the change.		

Prove each step before expanding AI.

First 90 days

Map eligible issues, record quality and cost baselines, and pilot invoice lookup and product how-to guidance. Decide on each pilot using lasting resolution, repeat contact, severe errors, and cases reaching people.

Year 1

Target 70% verified AI resolution and 66k weekly cases reaching people. Add coverage for the temporary queue increase while quality and backlog hold.

Year 2

Target 82% verified AI resolution and about 59k weekly cases reaching people. Show lower cost per lasting resolution while repeat contact and backlog stay at or below baseline.

Year 3

Test 90% verified AI resolution and 50k weekly cases reaching people. Adjust vendor capacity only after two consecutive reviews confirm strong quality, reliable coverage, and lower cost.

Check every year Report verified AI resolution, cases reaching people, lasting resolution, repeat contact, severe errors, backlog age, and cost per lasting resolution by issue type. The numeric targets above are my planning assumptions, separate from OpenAI commitments.

What I need from leadership

- Linked issue, repeat-contact, and cost data across channels.
- Named Product, Engineering, and policy owners.
- Training time and flexible partner terms.

How I would judge progress

- Durable resolution and severe errors show customer quality.
- Human case volume and backlog show capacity.
- Cost per lasting resolution shows the full economics.

Decision rule

Track one issue across AI, partners, and internal specialists, and count a lasting resolution once. Stop an affected AI action after a severe error or rising repeat contact. Keep human coverage if queues stay high.

The detailed owners, scorecard, pilot gates, and scenario math remain in the appendix for questions.

About \$118M in potential net cost avoided.

Compared with keeping AI resolution at 60% as demand rises. This is an illustrative forecast comparison, with no dollar credit for revenue. [Assumptions and sources.](#)

THREE-YEAR PLAN · ILLUSTRATIVE

~\$118M

Needs 90% vendor-eligible avoided cases, 3,000 cases per rep, and flexible contracts.

\$137.5M

Vendor spend avoided

\$20.0M

AI, tooling, and quality allowance

ONSHORE VENDOR REP

\$45/hour

About \$93.6k for a full year

OFFSHORE VENDOR REP

\$1,950/month

About \$23.4k for a full year

PLANNING RESULT	YEAR 1	YEAR 2	YEAR 3
Vendor spend avoided	\$6.10M	\$37.77M	\$93.63M
Program allowance	\$4.00M	\$6.00M	\$10.00M
Net impact	+\$2.10M	+\$31.77M	+\$83.63M

REWARD

Fund a staged pilot with case and contract data. If lasting resolution holds and Finance confirms vendor spend can be avoided, this plan reaches about \$118M net over three years.

RISK

At 82% AI resolution in Year 3, 70% vendor share, 4,000 cases per rep, and 10% budget conversion, the full \$20M spend leaves a \$13.1M shortfall. Stop expansion if those results emerge.

The math behind the stretch case.

PLANNING CASE	TODAY	YEAR 1	YEAR 2	YEAR 3
Weekly issues	150k	220k	330k	500k
Verified AI resolution	60%	70%	82%	90%
Cases reaching people	60k	66k	59k	50k

Starting point. The exercise supplies 1 billion weekly users, about 150k weekly issues, 60% automated resolution, and 40% human assistance.

Case complexity. As AI resolves more routine issues, I expect the remaining human queue to include more sensitive cases and exceptions. Partners and specialists would need training, current guidance, and clear escalation paths.


Staffing decision. I would use observed case volume, coverage, backlog, customer outcomes, and cost before adjusting staffing. The forecast assigns no headcount target.

Demand growth. User factors rise to 1.15×, 1.32×, and 1.50× across Years 1–3. Issues per user rise to about 1.28×, 1.67×, and 2.22× as products and support needs expand. These are aggressive scenario assumptions.

Slower automation. If verified AI resolution reaches only 82% in Year 3, about 90k issues would reach people each week, above today's 60k. I would keep coverage in place and investigate the issue mix before adjusting staffing.

Eligibility gate. I would measure eligible issue volume by type before endorsing the 90% target. The model gives no credit for product fixes that prevent an issue from arising.

How the model runs and recovers.



My approach Diagnose demand and case complexity, design the operating model, and deliver verified gains in stages.

Read the customer result and each contributor.

OPERATING LENS	MEASURE	LEADERSHIP RESPONSE
Quality	Durable resolution, repeat contact, audited answer quality, severe errors.	Stop the affected AI path if severe errors or repeat contact rise.
Productivity	Verified automation, human case volume, time to resolution.	Expand AI only when customer quality holds and fewer cases reach people.
Capacity	Qualified hours, backlog age, coverage by issue type, attrition, no-shows.	Move people to higher-skill tasks first and protect service coverage.
Cost to serve	Cost per durable resolution, including handoffs, rework, and transition costs.	Change vendor volume only when full cost improves.

One issue history shows each contribution		
AI	Partners	Internal teams
Eligible coverage, verified resolution, handoff quality, and error rate.	Lasting resolution, audit quality, staffing reliability, and rework on comparable cases.	Complex-case outcomes, escalation quality, and repeat demand removed through fixes.

Proposed pilots. Invoice lookup and product how-to guidance are examples. A Day 30 review would confirm eligible volume, risk, baseline, and the right review window.

Illustrative Day 90 gates. I would look for a 5-point gain in durable resolution, 10% fewer eligible cases reaching people, and zero severe errors before expansion. These are planning gates, separate from OpenAI commitments.

External source [OpenAI's published support model](#). Its support surfaces, knowledge, evaluations, and frontline feedback inform this proposal.

This is my independent recommendation for a hypothetical exercise. I have no access to OpenAI internal support data.

How the savings case is calculated.

INPUT	PLANNING VALUE	REASONING
Vendor rep cost	\$93.6k onshore \$23.4k offshore	Published rates of \$45 per hour and \$1,950 per month, carried across one full year.
Vendor case share	90%	AI removes repeatable cases that partners might otherwise handle. Issue-level data must confirm this share.
Vendor location mix	30% onshore 70% offshore	Illustrative mix for sensitive coverage and defined high-volume queues.
Cases per rep	3,000 each year	About \$14.82 per case at the blended vendor price. Case handling data must confirm throughput.
Budget conversion	40%, 75%, 90%	Contract changes or avoided future seats put potential value into the budget across Years 1-3.
Program allowance	\$4M, \$6M, \$10M	Illustrative AI usage, tooling, training, and oversight budget. Finance must price it.

Comparison. Keep automated resolution at 60% as weekly issues grow to 220k, 330k, and 500k. The proposed 70%, 82%, and 90% path avoids about 22k, 72.6k, and 150k human cases each week against that forecast.

Downside. With 65%, 75%, and 82% AI resolution, 70% vendor share, 4,000 cases per rep, and only 10% budget conversion, the full \$20M spend creates a \$13.1M shortfall. Staged funding limits that exposure.

Calculation. Annual cases avoided × 90% vendor share ÷ 3,000 cases per rep × \$44.46k blended annual price gives potential vendor capacity. Contract changes put \$137.5M into the three-year budget. After the \$20M program allowance, \$117.5M remains.

Scope. This measures expense avoided against the forecast. Today's vendor budget could still rise. Revenue, retention, product fixes, and internal staffing changes receive no dollar credit. Actual prices, case mix, AI cost, and quality remain unknown.

Public price references [SupportOps Global's published onshore and Philippines prices](#), [ConnectPro's Philippines rates](#), and [PITON-Global's fully loaded range](#). The [U.S. Bureau of Labor Statistics wage figure](#) provides a wage cross-check. Vendor bills include additional costs.

What changes if growth is slower.

This alternative reaches 281k weekly issues and 88% verified AI resolution in Year 3. It uses 6,000 cases per vendor rep and a \$9M program allowance. The public vendor prices and 30% onshore, 70% offshore mix stay in place.

ALTERNATIVE PATH	YEAR 1	YEAR 2	YEAR 3
Weekly issues	186k	230k	281k
Verified AI resolution	70%	80%	88%
Cases reaching people	56k	46k	34k
Net with 70% vendor share and slower contracts	−\$1.5M	+\$2.0M	+\$9.8M
Net with 90% vendor share and faster contracts	+\$0.58M	+\$8.96M	+\$20.56M

More cautious case. If 70% of avoided cases are vendor eligible and 10%, 40%, and 65% of capacity value reach the budget, the three-year net is \$10.3M.

Operating implication. Lower demand makes it easier for human cases to fall from 60k today to about 34k in Year 3. The financial opportunity is smaller because fewer future vendor seats need to be avoided.

Faster contract case. If 90% of avoided cases are vendor eligible and 40%, 75%, and 90% reach the budget, the three-year net is \$30.1M.

Decision rule. Use observed issue arrivals, vendor case time, contract terms, customer outcomes, and AI cost to move between these paths. Both are planning scenarios, separate from OpenAI targets.

Presentation choice The core pages use the higher-demand plan because it directly tests Support at hypergrowth. This lower-demand sensitivity gives leadership a conservative comparison for funding and staffing decisions.