

Build capacity through verified resolution.

My recommendation is to make every support path accountable for a durable customer outcome. Improve product help and AI for suitable issues, use partners for repeatable work, and reserve internal specialists for complex or sensitive cases. One operating team should own forecasting, quality, partner governance, and the feedback loop into Product and Engineering.

THREE-YEAR HORIZON

15–20 MINUTE READOUT

ALL EXERCISE DATA HYPOTHETICAL

THE DECISION

Fund the operating system before scaling the channel mix.

A higher automated-resolution rate helps only if users stay resolved and complex cases reach a capable owner.

- 01 Measure the full customer journey, including repeat contact.
- 02 Scale AI and partner capacity behind quality gates.
- 03 Turn recurring issues into product and knowledge fixes.

Volume grows. Human effort grows faster.

The assignment starts at one billion weekly users and 150,000 weekly support issues. At the stated 60% automated resolution rate, 90,000 issues finish in automation and 60,000 reach people. These exercise inputs are separate from reported OpenAI results.

TODAY / EXERCISE BASELINE

150k

weekly issues



90k automated · 60k human-assisted

YEAR 3 / HOLD CURRENT RATE

281k

weekly issues



169k automated · 112.5k human-assisted

146k complexity-adjusted human units

YEAR 3 / PROPOSED CASE

281k

weekly issues



211k automated · 70.3k human-assisted

91k complexity-adjusted human units

EXPLICIT PLANNING ASSUMPTIONS

Year-three users grow 1.5×. Weekly issues per user rise 1.25× as products and customer needs expand. The average human case needs 1.3× today's effort. A 75% verified automated-resolution rate is a conditional goal, subject to audits and user outcomes.

THE MATH

150k × 1.5 × 1.25 = 281k issues

281k × (1 - 0.75) × 1.3 ≈ 91k human effort units



The capacity implication: even if automation reaches 75%, human effort is about 52% above today's 60k baseline. If the automated rate stays at 60%, effort is about 144% above baseline. Plan people and partners for both paths until data narrows the range.

Three investments, in dependency order.

Begin with a reliable baseline, then improve resolution paths and partner execution. Automated share should follow better outcomes.

01

Instrument the journey

Create a shared issue taxonomy across product, channel, region, customer segment, partner, and risk tier. Link automation attempts, human handoffs, repeat contacts, and final outcomes to one issue ID.

EXPECTED EFFECT

Expose false closures and give workforce planning a case-mix baseline. Reported containment can hide unresolved users.

02

Improve the resolution system

Place help inside the product for common issues. Maintain versioned knowledge, test AI on representative cases, and let agents submit failed responses for review. Pilot low-risk actions behind approval and rollback controls.

EXPECTED EFFECT + TRADEOFF

Improve durable resolution and reduce repeat demand. Ship fewer automations while evaluation coverage catches up.

03

Rebuild partner capacity

Run one vendor-management system across sites. Route stable cases to trained partners, cross-train for peaks, and protect specialist capacity for product defects, fraud, account risk, and escalations.


EXPECTED EFFECT + TRADEOFF

Reduce backlog swings and improve partner quality. Carry reserve capacity for launches and incidents.

OpenAI has described support as a system connecting customer surfaces, knowledge, evaluations, and frontline feedback. The proposed investments extend that published direction to the exercise's capacity and partner decisions. [OpenAI support model](#) ↗

One journey. Distinct owners.

Route by issue risk and required skill, while keeping one customer record and one set of outcome definitions.
Support Operations owns the handoffs between teams.



Keep these decisions inside OpenAI

- Routing policy, risk thresholds, knowledge ownership, and AI release gates.
- Forecasts, capacity allocation, quality calibration, and partner contracts.
- Specialist escalation, incident command, and Product and Engineering priorities.

Ask partners to own execution

- Staff and coach approved queues across regions and time zones.
- Reach audited quality and staffing commitments by case type.
- Supply root-cause examples, failed-answer flags, and tested knowledge suggestions.

PARTNER COMMERCIAL DESIGN

Pay for dependable capacity and verified outcomes.

Use a base capacity commitment with variable rewards for audited resolution and quality. Compare partners within equivalent case and risk mixes. Hold a weekly service review, a monthly cross-partner calibration, and a quarterly commercial review. Maintain a trained backup route for critical queues and test the transfer plan each quarter. Decide on supplier consolidation only after the baseline exposes overlap and risk.

ISO 18295 covers contact-centre service and client oversight. COPC's customer-operations framework links forecasts, quality, knowledge, supplier management, and cost. The UK Sourcing Playbook supports outcome-focused measures, phased pilots, and supplier continuity plans. [ISO](#) [COPC](#) [Sourcing Playbook](#)

Earn the right to scale each change.

Each phase produces evidence for the next investment. The plan can accelerate when pilots pass and slow when customer outcomes fall.

FIRST 90 DAYS

Find the baseline

Audit the 60% automated-resolution claim. Segment weekly issues by intent, risk, channel, region, and repeat contact. Build a weekly demand forecast and shared partner scorecard. Select two high-volume, low-risk journeys for pilots.

EXIT EVIDENCE

Issue IDs join across channels, QA calibrates across teams, and pilots have comparison cohorts.

YEAR 1

Prove the system

Deploy versioned knowledge and AI evaluation. Pilot in-product help, agent assistance, and one approved action path. Rebalance partner queues after quality and staffing reviews. Establish specialist escalation and incident runbooks.

EXIT EVIDENCE

Durable resolution improves without a higher severe-error rate or longer specialist backlog.

YEAR 2

Expand by journey

Extend successful patterns across languages and product surfaces. Cross-train partner sites, test backup routing, and use frontline defect clusters to prioritize Product and Engineering fixes. Renegotiate contracts around measured outcomes.

EXIT EVIDENCE

Case-mix-adjusted partner performance stays stable through launches and regional peaks.

YEAR 3

Operate adaptively

Refresh forecasts monthly and after major launches. Route by skill and risk signals. Retire weak automations, expand tested ones, and allocate specialist capacity to emerging products and sensitive cases.

EXIT EVIDENCE

Service quality and backlog stay inside agreed limits at observed demand.

Score the journey, then its contributors.

Leadership should review outcomes by customer segment, issue type, region, and risk. Give operational credit after durable resolution, with each team’s contribution visible in the linked issue record.

MEASURE	DEFINITION AND OWNER	DECISION IT INFORMS
Durable resolution NORTH STAR	Issues resolved with no linked recontact within seven days, supported by sampled accuracy and customer feedback. Support Operations owns the definition.	If this falls, stop optimizing automated share or handle time until the failure is located.
Verified automated resolution	Eligible issues resolved by automation that pass the seven-day and sampled-quality checks. Automation Engineering owns changes; Support Operations audits.	Expand only in issue types whose quality reaches the approved pilot baseline.
Human service health	Useful first response, age of open issues, specialist handoff time, audited accuracy, and reopen rate. Review by case and risk mix.	Shift trained capacity when older cases or severe queues breach agreed limits.
Partner contribution	Case-mix-adjusted durable resolution, audited quality, staffing adherence, cost per durable resolution, and accepted knowledge fixes.	Coach, rebalance, or change partner scope based on repeated gaps within partner control.
Internal specialist contribution	Complex-case accuracy, escalation completion, defect confirmation, and knowledge fixes, adjusted for case severity. Specialist leads own review.	Add expert capacity or fix handoffs when severe cases age or bounce between teams.
End-to-end cost	Fully loaded AI, partner, internal, quality, and rework cost divided by durable resolutions. Finance and Support Operations reconcile the inputs.	Invest where quality-adjusted cost falls without moving effort or errors into another channel.
Product learning	Recurring-issue rate, confirmed defect time to fix, knowledge release time, and issue-rate change after a controlled product fix.	Fund product fixes when controlled data shows a stronger reduction in repeated demand.

PAUSE

Any severe privacy, fraud, policy, or account-access error pauses the affected route and sends the case to its human owner.

INVESTIGATE

Two weekly cohorts with recontact above the comparison cohort or accuracy below it stop expansion while the team reviews cases.

EXPAND

Improved durable resolution with stable risk and backlog qualifies a journey for staged expansion and continued sampling.

A study of 5,179 support agents found a 14% average productivity increase with AI assistance, with larger gains for newer or lower-skilled agents. This supports measured coaching pilots. Keep capacity gains outside this forecast until measured. NIST guidance supports human oversight, monitoring, and recovery paths for AI systems. [NBER study ↗](#) [NIST AI RMF ↗](#)

The operating pattern matches my experience.

The assignment needs a leader who can govern partners, coordinate across markets, and put AI into workflows with clear human ownership.

META / PARTNER GOVERNANCE

Five BPO partners across onshore and offshore programs.

For four years I managed partner performance across Meta programs. That experience informs the scorecard, staffing reviews, and recovery path I recommend here. The teams were vendor-employed, with my direct managers overseeing them.

GOOGLE / CROSS-MARKET SUPPORT

Product strategy and activation for Japan and Korea.

I worked with product and sales teams while supporting both markets. That experience informs regional handoffs, launch feedback, and clear ownership when one product serves different customer contexts.

ROWLAND AI / WORKFLOW DESIGN

AI workflows with evaluations and human gates.

At Rowland AI, I build agent workflows with tool integrations, evaluations, and guardrails. I would use that discipline to pilot support assistance, measure errors, and keep sensitive decisions with accountable people.

MY CLOSING
RECOMMENDATION

I would approve a 90-day baseline and two controlled journey pilots, then fund expansion against durable resolution, risk, and capacity evidence. That gives Support a way to grow with the product while showing leadership which investments help customers and which ones need correction.

Assumptions and source trail.

The recommendation stands on its own above. These notes expose the planning math and external evidence for asynchronous review.

Planning math

The exercise supplies 1B weekly users, 150k weekly issues, and 60% automated resolution. Year-three 1.5× user growth, 1.25× issue propensity, 1.3× human case effort, and 75% verified automated resolution are planning assumptions. They are sensitivity inputs, separate from forecasts or OpenAI targets.

At a constant 60% automated rate, $281,250 \times 40\% \times 1.3 = 146,250$ effort units. At 75%, $281,250 \times 25\% \times 1.3 = 91,406$ units. One unit equals today’s average human case effort. Staffing requires observed handle time, occupancy, shrinkage, service windows, and case mix before an FTE estimate.

In a stress case with 65% automated resolution and 1.5× human case effort, human effort reaches about 148k units. Maintain a backup capacity path until data narrows this exposure. Seven-day recontact is an initial measure. Longer case types need a longer outcome window.

Research used

- [OpenAI’s published support model](#) connects surfaces, knowledge, evaluations, and frontline feedback.
- [COPC CX Standard](#) informs service journeys, digital channels, quality, and vendor management.
- [ISO 18295-2](#) covers client oversight of in-house and outsourced contact centres.
- [UK Sourcing Playbook](#) informs outcome measures, pilots, and continuity planning.
- [Brynjolfsson, Li, and Raymond](#) provide support-agent AI assistance evidence.
- [NIST AI Risk Management Framework](#) informs oversight and monitoring.