

# Support at 1 billion weekly users

The exercise starts with 150,000 issues each week. Automation resolves 60%, leaving 60,000 issues for people. Growth and a harder case mix will raise human effort even if the number of human cases grows slowly.

|                                              |                                          |                                                |                                                 |
|----------------------------------------------|------------------------------------------|------------------------------------------------|-------------------------------------------------|
| <div>1 billion</div> <div>weekly users</div> | <div>150k</div> <div>weekly issues</div> | <div>60%</div> <div>automated resolution</div> | <div>60k</div> <div>human-assisted issues</div> |
|----------------------------------------------|------------------------------------------|------------------------------------------------|-------------------------------------------------|

## The decision for leadership

Invest in a support system that resolves more issues where customers already are, keeps skilled people available for harder cases, and turns repeat failures into product fixes.

The four starting figures come from the hypothetical take-home exercise. My growth and improvement figures below are scenario assumptions.

# Diagnose · Design · Deliver

Diagnose the full customer journey and the capacity left after automation. Design a shared system for product help, partner teams, and internal specialists. Deliver in stages, expanding only when customer outcomes and operating capacity support the change.

01

**Diagnose**

Connect the issue, channel, answer, handoff, repeat contact, and final customer outcome. Forecast effort by case type instead of using ticket counts alone.

02

**Design**

Put tested help inside the product, give partners clear queues and quality rules, and reserve internal specialists for judgment and product defects.

03

**Deliver**

Pilot two common journeys, review matched outcomes, train backup coverage, and move recurring failure patterns into Product and Engineering.

**The vision**

What this looks like in 3 years

A customer gets useful help in the product, and the issue history follows them if a person steps in. Support teams, partners, and product owners use one view of quality, effort, and repeat demand. Specialists spend more time on complex cases and on improving the system.

**Where to play**

Start with high-volume, low-risk journeys where the customer need and outcome can be measured. Keep sensitive, unclear, or high-risk issues with trained people.

**How to win**

Connect every support step to the customer outcome, then improve the knowledge, routing, tooling, and partner coaching behind it.

# Human cases level off while effort keeps rising.

I use a three-year planning case to show the capacity implications. By Year 3, weekly users reach 1.5× today's exercise baseline, issues per user reach 1.25×, and a human case takes 1.3× today's average effort. Verified automated resolution rises from 60% to 75% in stages.

**Why these assumptions:** User growth expands the base. More products and business use cases can lift issue frequency. Automation can remove easier cases first, leaving a harder human queue. The multipliers illustrate those pressures and a measured improvement path; they come from my judgment, separate from any OpenAI forecast.

| PLANNING VIEW                                                   | WEEKLY ISSUES | VERIFIED AUTOMATION | HUMAN CASES | HUMAN EFFORT UNITS |
|-----------------------------------------------------------------|---------------|---------------------|-------------|--------------------|
| <b>Today</b><br>1.00× users · 1.00× issues/user · 1.00× effort  | 150k          | 60%                 | 60k         | 60k                |
| <b>Year 1</b><br>1.15× users · 1.08× issues/user · 1.10× effort | 186k          | 65%                 | 65k         | 72k                |
| <b>Year 2</b><br>1.32× users · 1.16× issues/user · 1.20× effort | 230k          | 70%                 | 69k         | 83k                |
| <b>Year 3</b><br>1.50× users · 1.25× issues/user · 1.30× effort | 281k          | 75%                 | 70k         | 91k                |

Effort units use today's average human case as 1 unit. I have rounded the figures. If automation stays at 60% in Year 3, the model reaches about 113k human cases and 146k effort units.

### Capacity and capability trend

Human case volume rises modestly in this planning case, while effort climbs about 52% from today's baseline. I would add specialized skills, better tools, and trained flexible coverage before translating effort units into staffing. Forecasting would use handle time by case type, qualified hours, occupancy, shrinkage, and service windows.

# Three investments make the model usable.

|                                                                                                                                                                                                                                                                                               |                                                                                                                                                                                                                                                                                                              |                                                                                                                                                                                                                                                                                                    |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <div>01</div> <div>One issue record</div> <div>Join channel, automation, partner, specialist, and repeat-contact events. This makes customer outcomes and contribution visible by issue type.</div> <div>Trade-off</div> <div>Data integration comes before broad automation expansion.</div> | <div>02</div> <div>Knowledge and evaluations</div> <div>Turn frontline failures into updated guidance, test cases, and product fixes. Roll out low-risk journeys with sampled review and an easy human handoff.</div> <div>Trade-off</div> <div>Controlled pilots grow slower, with stronger evidence.</div> | <div>03</div> <div>Partner capability</div> <div>Give partners clear routing, calibrated quality, training, and trained backup staffing. Shift incentives toward durable resolution and lower rework.</div> <div>Trade-off</div> <div>Coverage and training cost precede productivity gains.</div> |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

## Who does what

|                                                                                                                                                                                                                                                                       |                                                                                                                                                                                                                                                                 |                                                                                                                                                                                                                                                                                                                          |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <div>Product help and AI</div> <div>Handles approved, repeatable requests within clear limits.</div> <div>Owns a complete history, a visible handoff route, and traceable answers.</div> <div>Improves through evaluations, content changes, and product fixes.</div> | <div>Partner teams</div> <div>Handle repeatable human cases across approved queues and time zones.</div> <div>Own staffing, coaching, documentation, and trained backup coverage.</div> <div>Improve through audited outcomes and issue-pattern feedback.</div> | <div>Internal specialists</div> <div>Handle sensitive issues, unclear policy, incidents, and product defects.</div> <div>Own routing rules, quality calibration, forecasts, partner governance, and final escalation decisions.</div> <div>Improve the knowledge and product backlog with Product and Engineering.</div> |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

OpenAI would set policy, evaluation rules, launch gates, and partner contracts. Partners would execute within those rules. A severe privacy, fraud, policy, or access error would route to a human owner and stop the affected automated path until reviewed.

# Judge the full customer outcome, then the contributors.

The lead measure is durable resolution by issue type. The customer gets a correct answer with no repeat contact about that issue within the appropriate review window. A shared issue ID lets leadership see what automation, partners, and internal specialists each contributed.

| QUESTION                            | MEASURES TO REVIEW                                                                                                                                        | WHAT WOULD CHANGE MY PLAN                                                                       |
|-------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------------------|
| Did the customer get help?          | Durable resolution, repeat contact, audited answer quality, severe errors, time to resolution, and customer effort.                                       | Rising recontact or severe errors would stop expansion on that route and trigger case review.   |
| Can capacity keep up?               | Arrivals by issue type, qualified hours, backlog age, occupancy, handle time, partner attrition, no-show rate, and late rate.                             | Persistent aged backlog would shift schedules, trained backup staffing, or specialist capacity. |
| Does each step add value?           | Eligible automation use, verified automation resolution, handoff quality, partner durable resolution, specialist rework, and cost per durable resolution. | Higher automation use with weaker outcomes would send the journey back to design.               |
| Where is repeat demand coming from? | Issue coverage across products and channels, journey starts, successful routing, resolution, and repeat-contact clusters.                                 | Recurring clusters would move into a Product and Engineering fix queue.                         |

## How I would adapt my Meta scorecard

**Revenue.** Gross and incremental revenue help prioritize business customer pain, while durable resolution remains the Support goal.

**Solution score.** Pitched calls per book and solutions per call become the share of eligible issues with a relevant next step and a verified outcome.

**Workforce.** Attrition, no-shows, and late arrivals still show staffing reliability. Very short contacts prompt a quality sample.

**Tooling.** Ads AI adoption becomes support-tool use among eligible agents, paired with audited quality and customer outcomes.

**Coverage.** Breadth becomes product, channel, and issue coverage. Depth becomes qualified staffing for each queue.

**Quality.** Tailored solutions and commercial conversations become answer relevance, context use, and clear ownership for business customers.

**Journey funnel.** Attempted, connected, pitched, and actioned become issue start, routing, guidance delivered, and confirmed resolution. Segment by entry channel.

Compare like-for-like issue types, risk, channel, and customer groups. Show total customer outcomes plus automation, partner, and internal specialist contributions; a transfer or reopened issue should remain visible across all three.

# Build the evidence, then expand the useful paths.

|               |                                                                                                                                          |
|---------------|------------------------------------------------------------------------------------------------------------------------------------------|
| First 90 days | Join issue and outcome data, set the quality and capacity baseline, choose two low-risk journeys, and pilot with human review.           |
| Year 1        | Improve product help and partner routing on proven journeys. Calibrate quality across partners and build trained backup coverage.        |
| Year 2        | Expand proven paths across more products and channels. Move recurring case clusters into Product and Engineering roadmaps.               |
| Year 3        | Operate one linked support system, with specialists focused on complex cases and partners handling well-defined queues at greater scale. |

● Stop

Stop a route after a severe privacy, fraud, policy, or access error.

● Start

Start a small pilot when its issue type, owner, outcome, and human route are defined.

● Continue

Continue staged expansion when durable resolution improves and risk and backlog stay controlled.

## Action plan

Status colors show illustrative Day 0 readiness for this exercise. They describe a proposed plan before execution begins. Red means a blocking input, yellow means design needed, and green means ready to start.

| ACTION                                                         | OWNER ROLE                      | TIMELINE   | ILLUSTRATIVE STATUS       |
|----------------------------------------------------------------|---------------------------------|------------|---------------------------|
| Connect issue history, automation results, and repeat contact. | Support Ops + Analytics         | Days 0-30  | ● Red · data access       |
| Agree on issue types, quality rules, and partner scorecards.   | Support Ops + Partner Lead      | Days 0-60  | ● Yellow · design         |
| Pilot two low-risk journeys with human handoff.                | Product + Engineering + Support | Days 30-90 | ● Yellow · design         |
| Map trained backup staffing and specialist escalation.         | Workforce Planning + Partners   | Year 1     | ● Green · ready           |
| Expand proven journeys and fix repeat issue clusters.          | Support + Product + Engineering | Years 2-3  | ● Yellow · pilot evidence |

My close

I would ask for linked issue data, two pilot journeys, a Product and Engineering partner, and authority to set common partner quality rules. In 90 days, I would bring leadership an observed capacity plan, a decision on each pilot, and the next set of product fixes.

# The planning assumptions behind this recommendation.

**Starting point.** The exercise supplies 1 billion weekly users, about 150k weekly issues, 60% automated resolution, and 40% human assistance.

**Case complexity.** Human effort per case rises to 1.10×, 1.20×, and 1.30× because repeatable issues may leave the human queue first. One effort unit equals today's average human case.

**Capacity conversion.** The model estimates effort before translating it into headcount. Staffing requires handle time, qualified hours, occupancy, shrinkage, case mix, and service windows.

**Demand growth.** Weekly users rise to 1.15×, 1.32×, and 1.50× across Years 1–3. Issues per user rise to 1.08×, 1.16×, and 1.25×. These illustrative steps are my planning assumptions.

**Automation path.** Verified resolution rises to 65%, 70%, and 75% only in this planning case. The exercise provides no target. Count an issue as resolved only after an appropriate repeat-contact window and quality check.

**Contribution and risk.** A linked issue record attributes each step while preserving the final customer outcome. Human owners review severe errors and expansion decisions.

**External source:** [OpenAI's published support model](#). Its discussion of support surfaces, knowledge, evaluations, and frontline feedback informs the operating design.

This is my independent recommendation for a hypothetical exercise. I have no access to OpenAI internal support data.